

Databases project

PokéSallianWorld 2022-2023 – Stage 3



List of members (name and email):

Valeria Ezquerro Rodriguez	valeria.ezquerro@students.salle.url.edu
Eugènia Pacheco i Ferrando	eugenia.pacheco@students.salle.url.edu
Sebastián Felix Gorga	sebastianfelix.gorga@students.salle.url.edu
Marc Soto José	m.soto@students.salle.url.edu

Date of finalization: 28/05/2023

Index

1 INTRODUCTION	4
2 ENTITY-RELATIONSHIP MODEL UPDATE	5
3 RELATIONAL MODEL UPDATE	6
4 PHYSICAL MODEL UPDATE	7
5 POKÉMON ARE UNIQUE WARRIORS	8
5.1 QUERY 1	8
5.2 QUERY 2	9
5.3 QUERY 3	11
5.4 QUERY 4	12
5.5 QUERY 5	14
5.6 QUERY 6	15
5.7 TRIGGER 1	16
5.8 TRIGGER 2	19
6 I'M NOT A GAMER, I'M A POKÉMON TRAINER	20
6.1 QUERY 1	20
6.2 QUERY 2	22
6.3 QUERY 3	23
6.4 QUERY 4	26
6.5 QUERY 5	28
6.6 QUERY 6	31
6.7 TRIGGER 1	33
6.8 TRIGGER 2	37
7 LET'S GO POKESHOPPING	41
7.1 QUERY 1	41
7.2 QUERY 2	42
7.3 QUERY 3	43
7.4 QUERY 4	45
7.5 QUERY 5	48
7.6 QUERY 6	49
7.7 TRIGGER 1	51
7.8 TRIGGER 2	53
8 TIME TO EXPLORE	58
8.1 QUERY 1	58
8.2 QUERY 2	60
8.3 QUERY 3	61
8.4 QUERY 4	63
8.5 QUERY 5	64
8.6 QUERY 6	66
8.7 TRIGGER 1	68
8.8 TRIGGER 2	71
9 CROSS QUERIES.	75
9.1 QUERY 1	75
9.2 QUERY 2	76
9.3 QUERY 3	77
9.4 QUERY 4	77
9.5 QUERY 5	80
9.6 QUERY 6	82

9.7 QUERY 7	82
9.8 QUERY 8	84
10 Conclusions	85
10.1 USE OF RESOURCES	85
10.2 IA USE (IF REQUIRED, 1-2 PAGES)	85
10.3 LESSONS LEARNT AND CONCLUSIONS (1 PAGE)	85

1 Introduction

In this last phase of the project we will validate the correct functioning of our database by performing several queries to retrieve relevant information about the game. As we demonstrated on the previous stage, we managed to import all the given data and populate our tables. Now it's time to manage them by taking specific information from particular tables while trying to optimize our queries as much as possible to ensure no unnecessary processing power is being used.

Additionally, we will automatize several procedures by creating *triggers* to guarantee the consistency of our data. That is, whenever some changes of data on a single table affect others, we will first validate they are in the correct format and make sense in the context of the table. We will also use them to detect suspicious activity and document it on a new **Warnings** table.

Similarly to previous stages, we have decided to split the work in modules (one for each team member). Once all the queries and triggers of each module are done and have been checked by everyone in the team, we will proceed to do the *cross queries*. These are the queries that require gathering information about tables belonging to different modules.

2 Entity-relationship model update

The only change we had to do was related to the **Cause_Or_Receive** table. On our database population we took the different moves provided in the **Pokemon_Instances.csv** file (*move1*, *move2*, *move3*, and *move4*) and we created new relationships using this table, but completely ignoring the number of this move (as we thought it wasn't relevant). However, we have now realized that it can be important to know which move is which, as in one of the triggers we have we must access the move that occupies the same slot as a given one. Therefore, we added a new attribute *slot* to this table and now when inserting the column *move1* onto our **Cause_Or_Receive** table, we also insert its *slot* as a 1 (same for the other three moves).

COLOR CODE

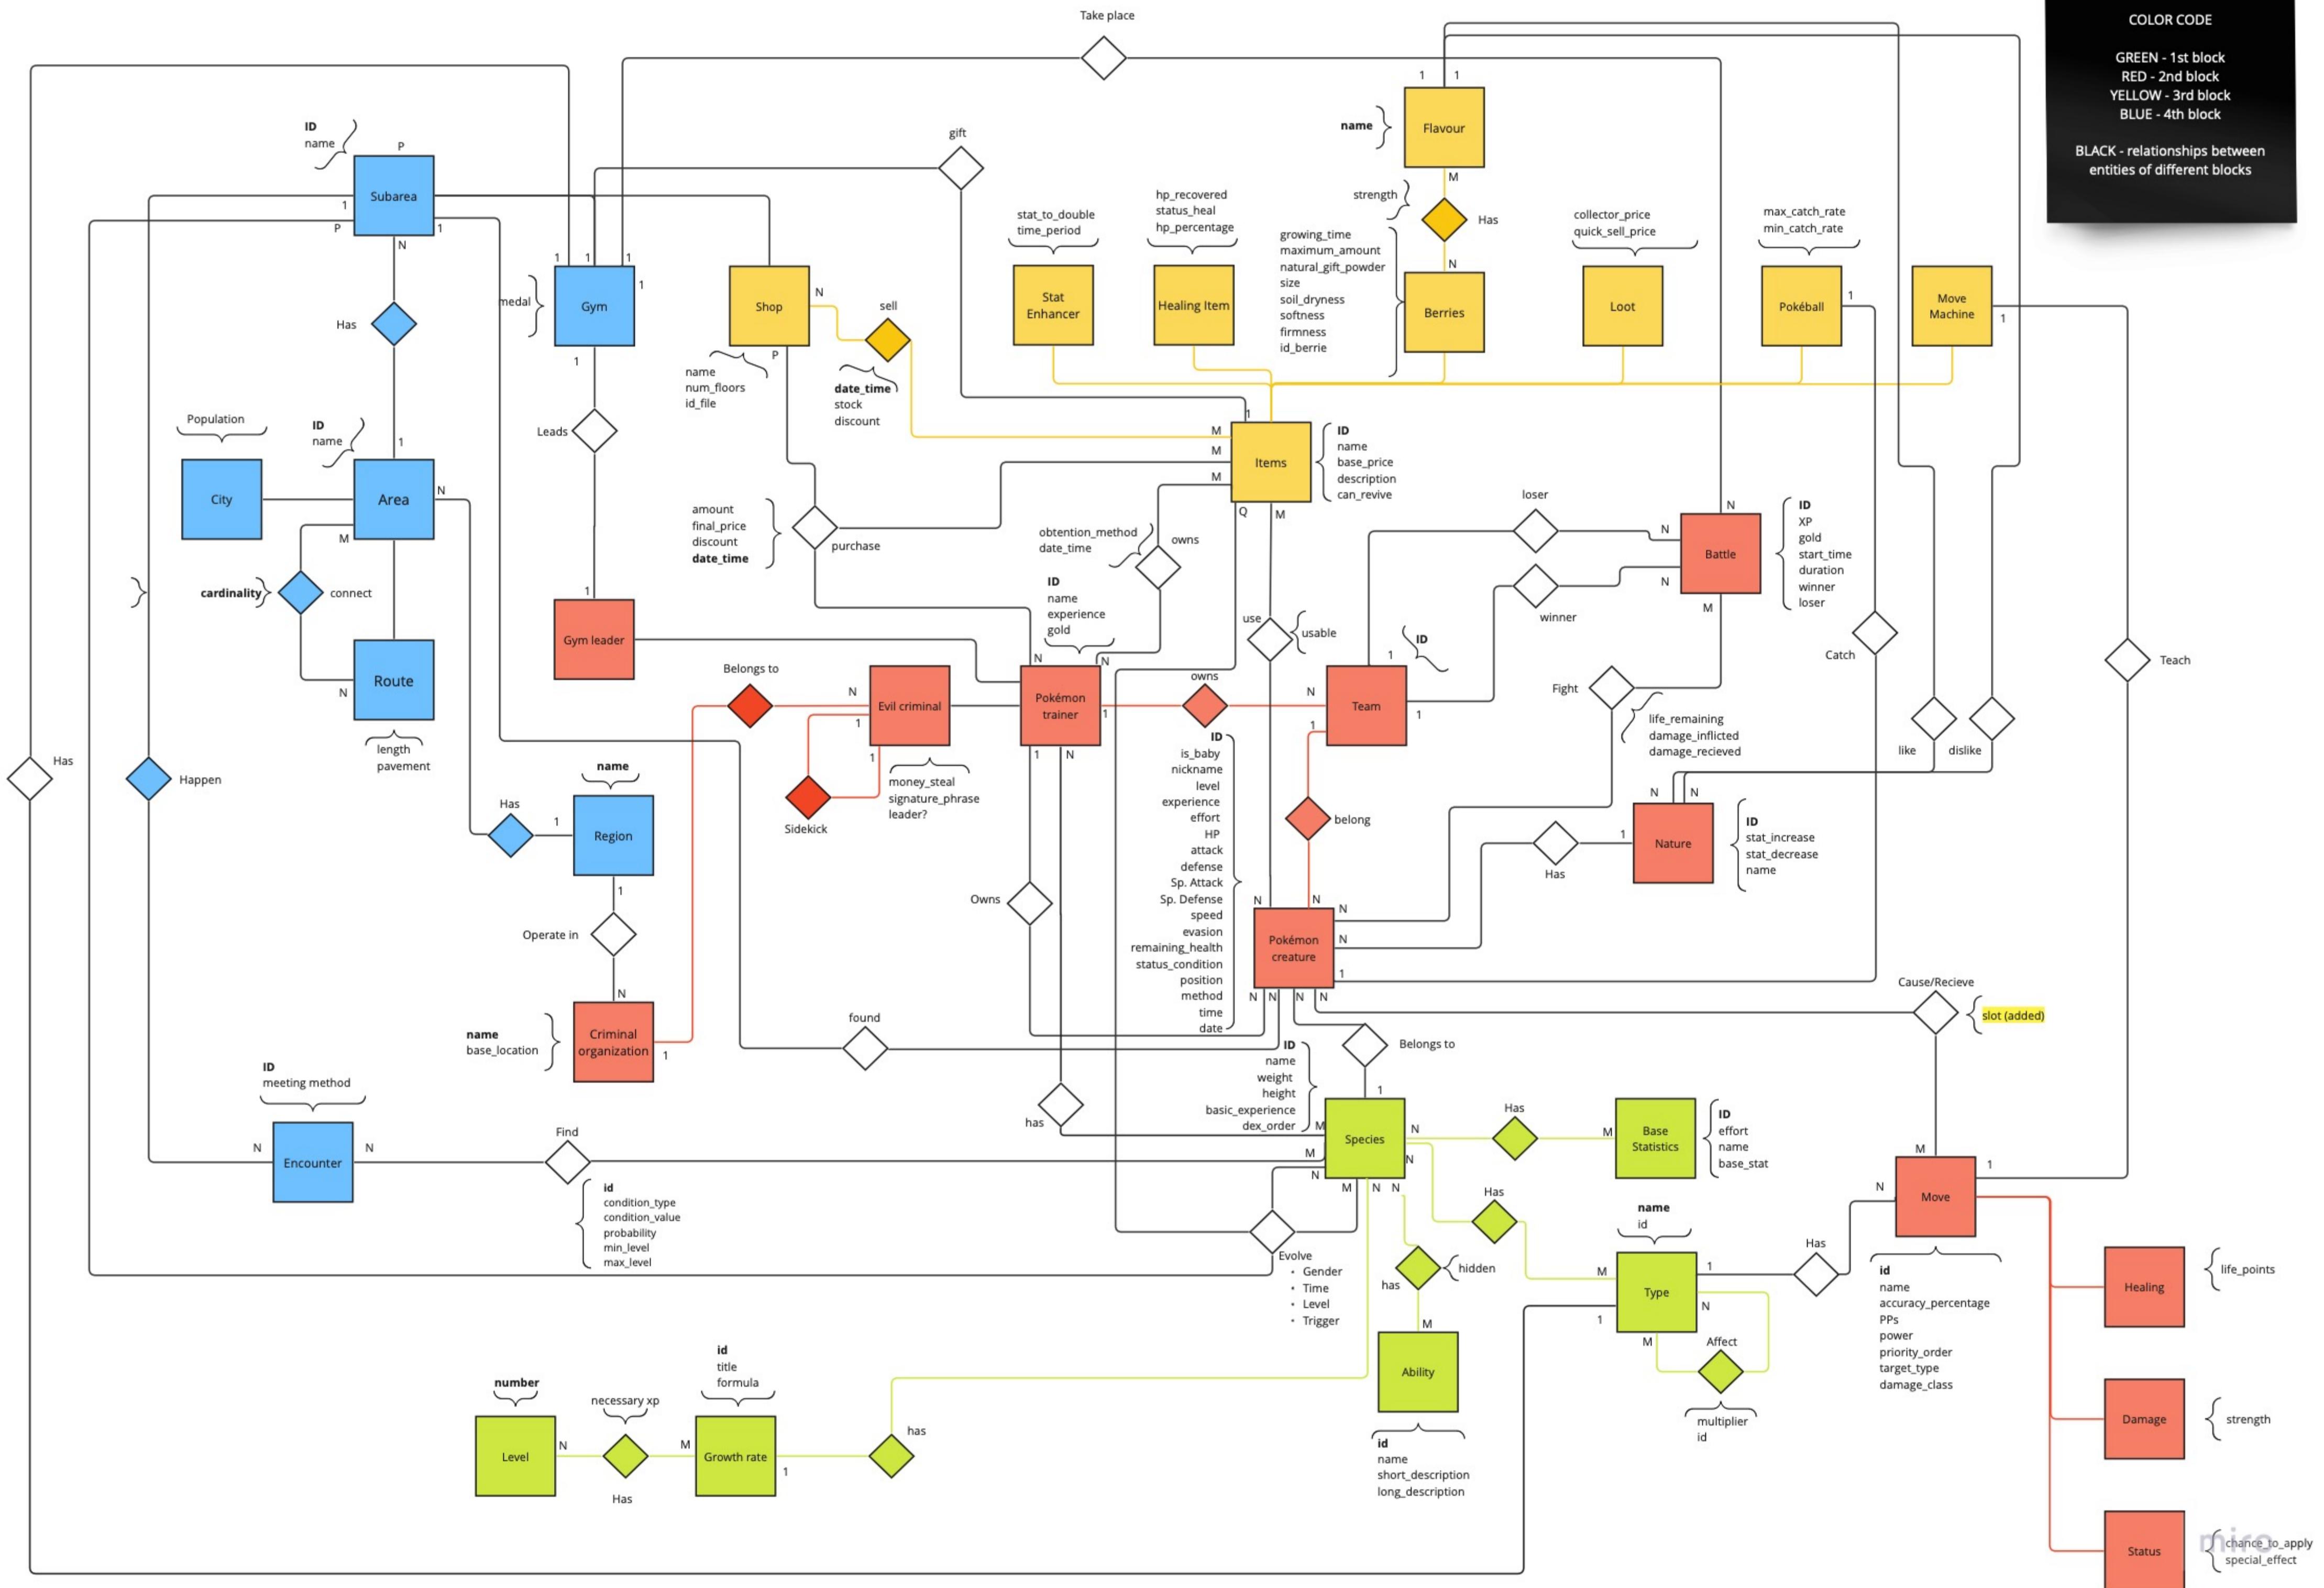
GREEN - 1st block

RED - 2nd block

YELLOW - 3rd block

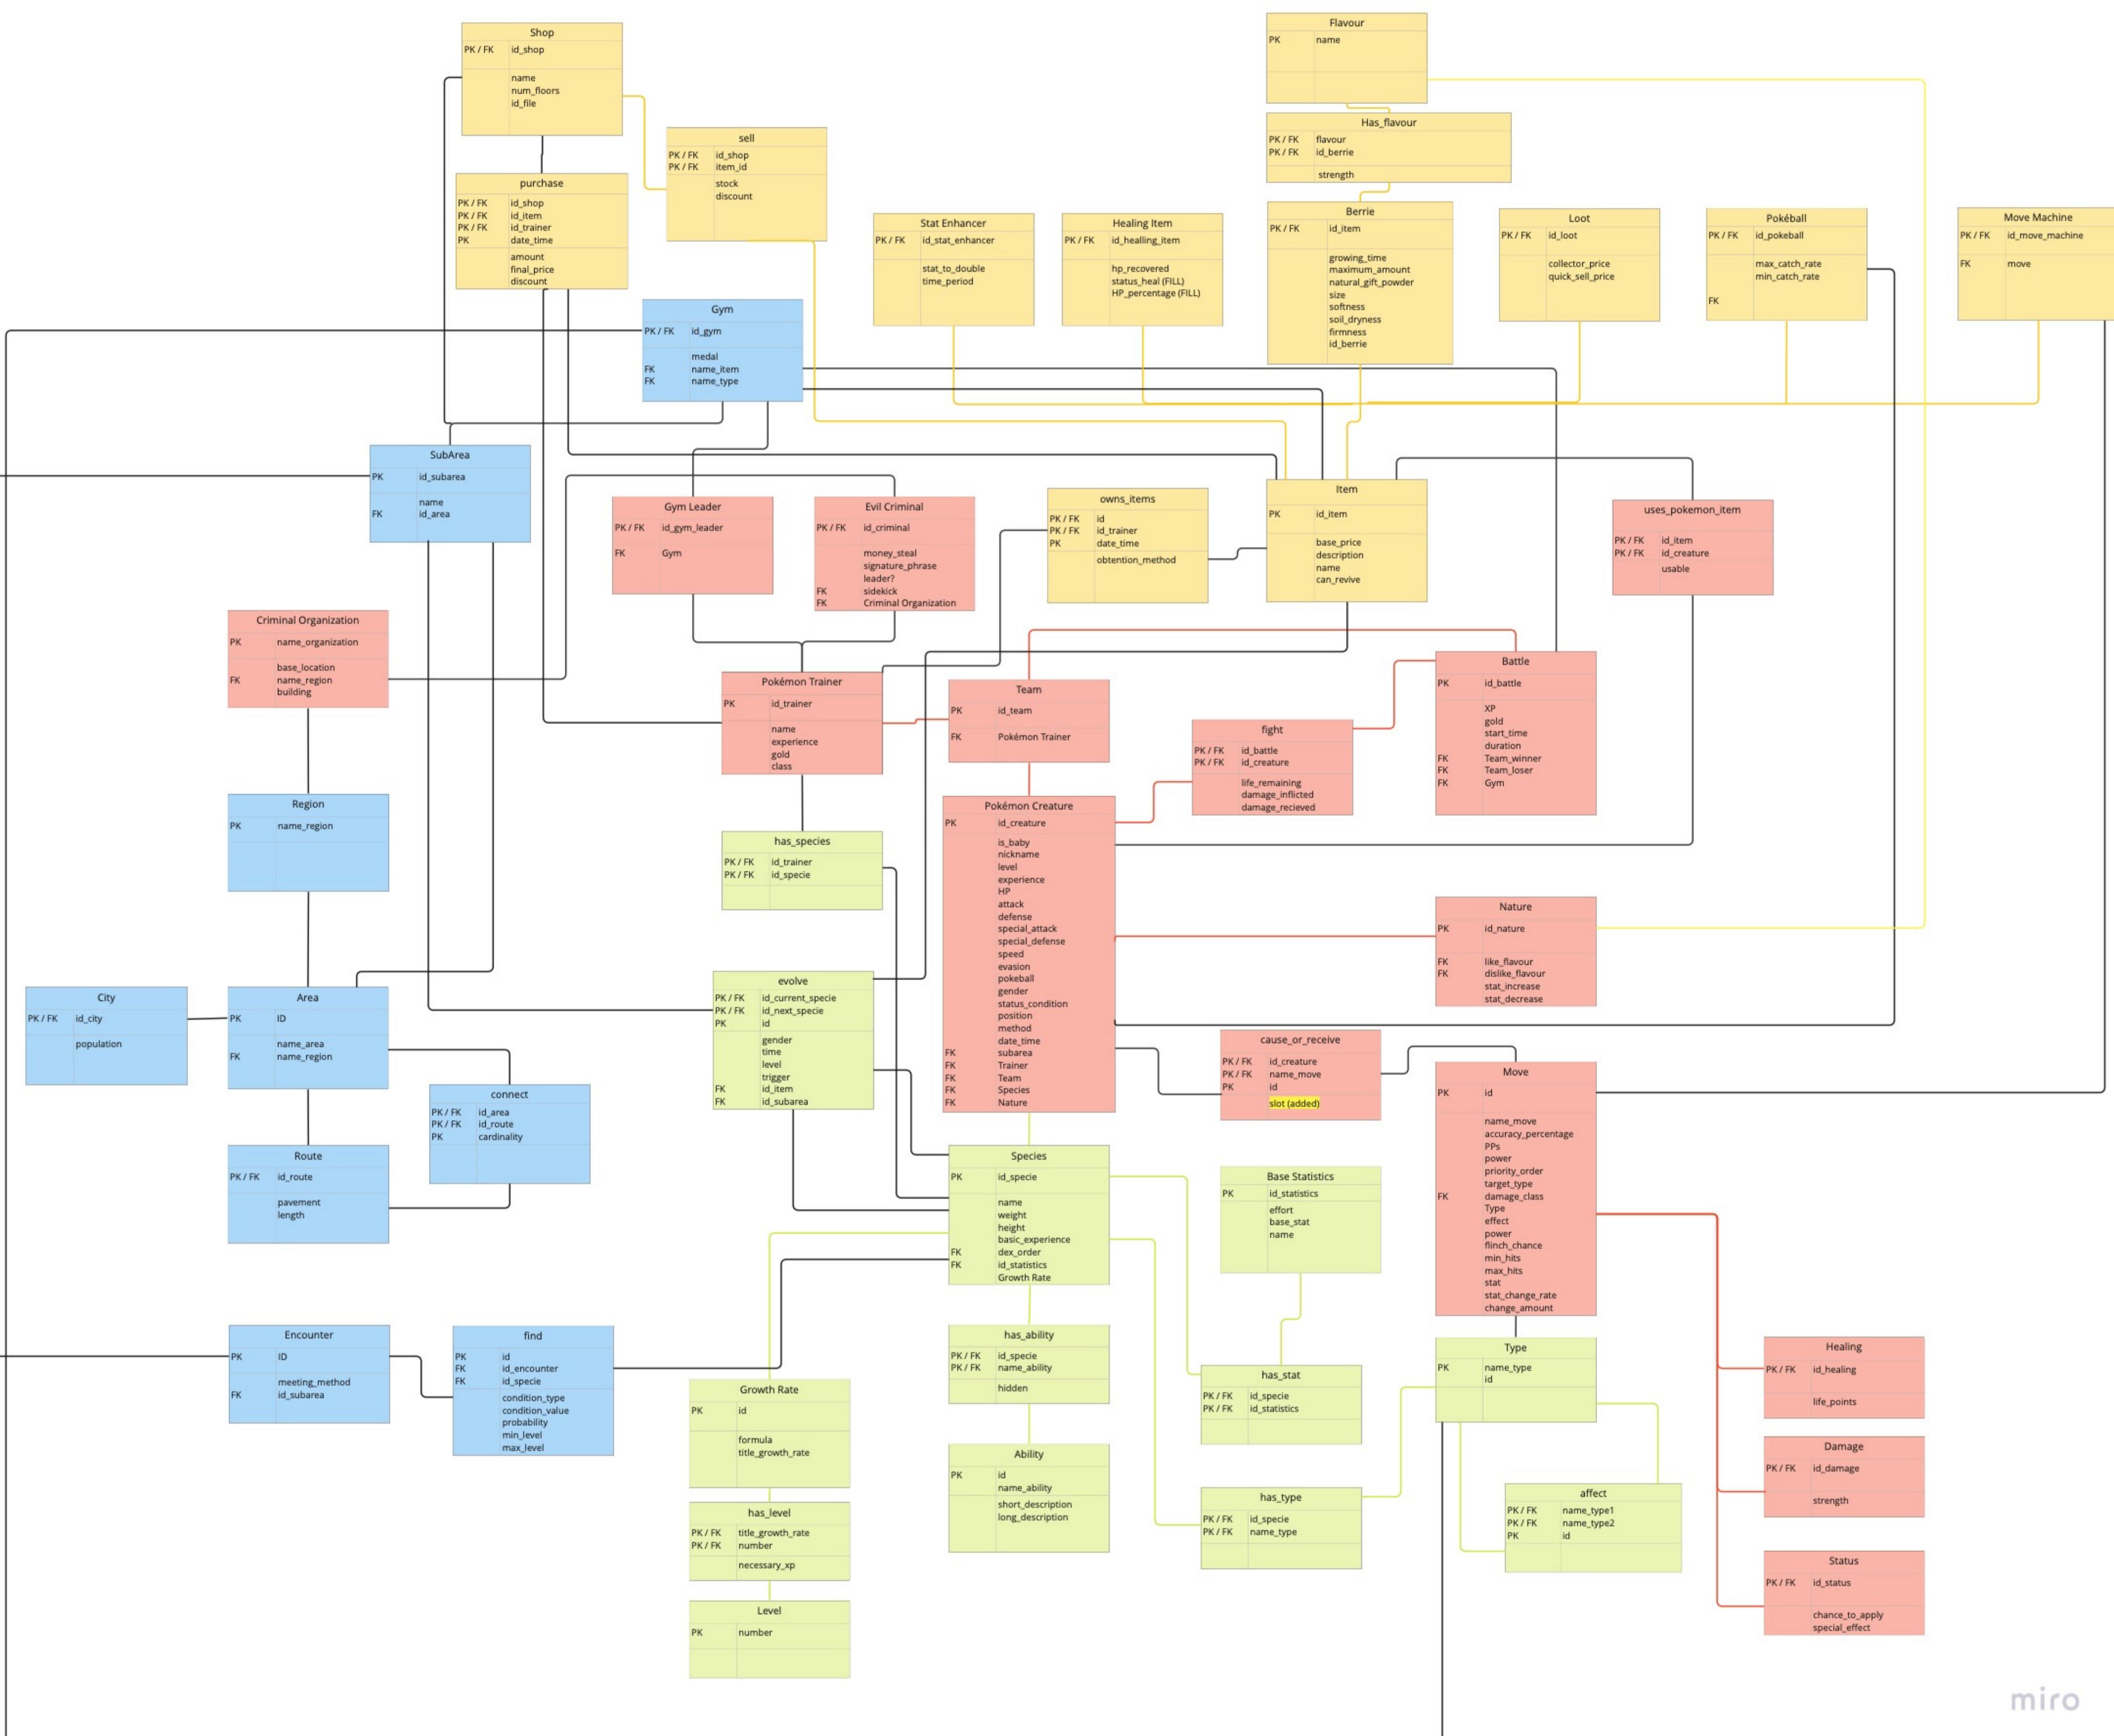
BLUE - 4th block

BLACK - relationships between entities of different blocks



3 Relational model update

As we mentioned in the previous section, we had to add a new attribute *slot* to the **Cause_Or_Recieve** table. We defined this one as a simple attribute, without declaring it as a FK or PK.



4 Physical model update

When trying to link the ***Criminal_Organization*** and ***Region*** tables in one of the queries we realized that the attribute *region* we had in this first table, hadn't been declared as a foreign key to the second. This was specified on our relational model diagram, but we had forgotten to add it to our physical model. We have fixed this little issue and now the query works as expected.

5 Pokémons are unique warriors

5.1 Query 1

5.1.1 Solution

```
SELECT DISTINCT s.name, bs.base_stat AS defense, bs2.base_stat AS special_defense, s.weight
FROM species AS s
JOIN has_stat AS hs1 ON s.id = hs1.id_species
JOIN base_statistics AS bs ON hs1.id_statistics = bs.id
JOIN has_stat AS hs2 ON s.id = hs2.id_species
JOIN base_statistics AS bs2 ON hs2.id_statistics = bs2.id
JOIN has_ability AS ha ON s.id = ha.id_species
JOIN ability a ON ha.id_ability = a.id
WHERE bs.name = 'defense' AND bs2.name = 'special defense' AND a.name = 'sturdy'
ORDER BY s.weight DESC
LIMIT 10;
```

	name	defense	special_defense	weight
1	steelix	200	65	4000
2	aggron	180	60	3600
3	probopass	145	150	3400
4	golem	130	65	3000
5	regirock	200	100	2300
6	onix	160	45	2100
7	magnezone	115	90	1800
8	bastiodon	168	138	1495
9	forretress	140	60	1258
10	donphan	120	60	1200

5.1.2 Explanation

As we want to output the *name*, the *defense* and *special defense* values together with the *weight*, we must select all of them. In order to get the desired data, we *JOIN* species with *has_stat* and *has_stat* with *base_statistics* twice since each species has 4 *base_statistics* but we just want 2 of them (*defense* and *special_defense*). After this, the *WHERE* clause is used to indicate that we want only those *defense* and *special_defense* values as long as the ability name is sturdy. Finally, we order the result by descendent weight and we limit to 10 as the statement says.

5.1.3 Query validation

In order to check this query, we get the id of the ‘sturdy’ ability and we filter by this ID in the *has_ability* table. This way, we obtain all the species IDs that have that particular ability. Then, by doing another query, we get all the species names ordered by descendent weight.

WHERE name = 'sturdy'	ORDER BY
1 5 sturdy	short_description long_description

```
SELECT ha.id_species, s.name
FROM species AS s
```

```

JOIN has_ability AS ha ON ha.id_species = s.id
WHERE id_ability = 5
ORDER BY s.weight DESC;

```

The screenshot shows two tables side-by-side. The left table is a result of a query with the following data:

	id_species	id_ability	hidden
1	74	5	false
2	75	5	false
3	76	5	false
4	81	5	false
5	82	5	false
6	95	5	false
7	185	5	false
8	204	5	false
9	205	5	false
10	208	5	false
11	213	5	false
12	227	5	false
13	232	5	false
14	299	5	false
15	304	5	false
16	305	5	false
17	306	5	false
18	369	5	true
19	377	5	true
20	410	5	false
21	411	5	false
22	438	5	false
23	462	5	false
24	476	5	false

The right table lists all species with their IDs and names:

	id_species	name
1	208	steelix
2	306	aggron
3	476	probopass
4	76	golem
5	377	regirock
6	95	onix
7	462	magnezone
8	411	bastiodon
9	205	forretress
10	232	donphan
11	305	lairon
12	75	graveler
13	299	nosepass
14	304	aron
15	82	magneton
16	410	shieldon
17	227	skarmory
18	185	sudowoodo
19	369	relicanth
20	213	shuckle
21	74	geodude
22	438	bonsly
23	204	pineco
24	81	magnemite

As we can see, the heaviest species with ‘sturdy’ ability match with our output.

Now that we have the heaviest species, we can check if the *defense* and *special_defense* values are correct for ‘steelix’ (for example). If we filter by its ID in the **has_stats** table, we get that the abilities it has are the following ones:

	id_species	id_statistics
1	208	1243
2	208	1244
3	208	1245
4	208	1246
5	208	1247
6	208	1248

Now, if we go to the table **base_statistics**, we can see that the values match with what we got.

	id	effort	base_stat	name
1	1245	2	200	defense
1	1247	0	65	special defense

5.2 Query 2

5.2.1 Solution

```

SELECT AVG(CASE WHEN (bs.name = 'attack' OR bs.name = 'special
attack') THEN bs.base_stat END) AS average_damage_bs,

```

```

        MAX(CASE WHEN (bs.name = 'attack' OR bs.name = 'special
attack') THEN bs.base_stat END) AS max_damage_bs,
        MIN(CASE WHEN (bs.name = 'attack' OR bs.name = 'special
attack') THEN bs.base_stat END) AS min_damage_bs
FROM species AS s
JOIN has_stat ON s.id = has_stat.id_species
JOIN base_statistics bs ON has_stat.id_statistics = bs.id
WHERE s.id IN (SELECT id_species
               FROM has_type
               WHERE type_name = 'fire');

```

	average_damage_bs	max_damage_bs	min_damage_bs
1	84.86363636363636	130	40

5.2.2 Explanation

In this query we want to get the *average*, *maximum* and *minimum damage base_statistic* of the species of *type* ‘fire’. Note that since we want to consider only those base stats related to damage, we must do a CASE when selecting the information. Then, we *JOIN species*, *has_stat* and *base_statistics* tables. In the *WHERE* clause, we state that we only want those species whose ID is included in *has_type* table and that the *type* is fire. To specify this, we use a subquery since we just want to take into account the species that have type ‘fire’.

5.2.3 Query validation

We can check that the MAX and MIN values are correct by doing the following query:

```

SELECT s.name, ht.id_species, MAX(CASE WHEN (bs.name = 'attack' OR
bs.name = 'special attack') THEN bs.base_stat END),
        MIN(CASE WHEN (bs.name = 'attack' OR bs.name = 'special
attack') THEN bs.base_stat END)
FROM species AS s
JOIN has_type AS ht ON s.id = ht.id_species
JOIN has_stat hs on s.id = hs.id_species
JOIN base_statistics bs on hs.id_statistics = bs.id
WHERE ht.type_name = 'fire'
GROUP BY(s.name, ht.id_species);

```

If we order the MAX stat from highest to lowest and the MIN stat from lowest to highest, we get this:

	name	id_species	max	min
1	flareon	136	130	95
2	ho-oh	250	130	110
3	heatran	485	130	90
4	magmortar	467	125	95

	name	id_species	max	min
1	slugma	218	70	40
2	vulpix	37	50	41
3	magcargo	219	90	50
4	charmander	4	60	52

As the screenshots show, the min base stat is 40 and the max base stat is 130.

As for the AVG value, by computing the following query we can see that it coincides:

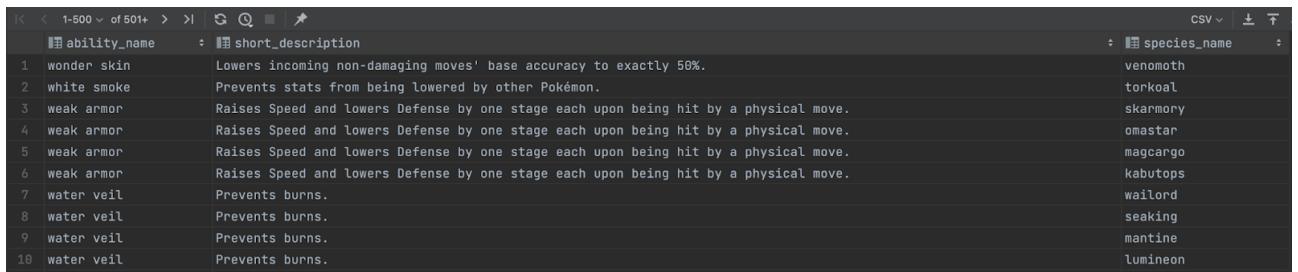
```
SELECT AVG(CASE WHEN (bs.name = 'attack' OR bs.name = 'special attack') THEN bs.base_stat END)
FROM species AS s
JOIN has_stat hs ON s.id = hs.id_species
JOIN base_statistics bs ON hs.id_statistics = bs.id
WHERE s.id IN (SELECT id_species
                FROM has_type
                WHERE type_name = 'fire');
```

	avg
1	84.86363636363636

5.3 Query 3

5.3.1 Solution

```
SELECT a.name AS ability_name, a.short_description, s.name AS species_name
FROM species AS s JOIN has_ability AS ha ON s.id = ha.id_species
JOIN ability AS a ON a.id = ha.id_ability
WHERE s.basic_experience > (SELECT AVG(basic_experience)
                             FROM species)
ORDER BY (a.name, s.name) DESC;
```



The screenshot shows a database interface displaying the results of a query. The results are presented in two columns: 'ability_name' and 'species_name'. The 'ability_name' column lists various abilities like 'wonder skin', 'white smoke', etc., with their descriptions. The 'species_name' column lists species names like 'venomoth', 'torkoal', etc. The interface includes navigation buttons at the top left and a CSV export option at the top right.

ability_name	short_description	species_name
wonder skin	Lowers incoming non-damaging moves' base accuracy to exactly 50%.	venomoth
white smoke	Prevents stats from being lowered by other Pokémon.	torkoal
weak armor	Raises Speed and lowers Defense by one stage each upon being hit by a physical move.	skarmory
weak armor	Raises Speed and lowers Defense by one stage each upon being hit by a physical move.	omastar
weak armor	Raises Speed and lowers Defense by one stage each upon being hit by a physical move.	magcargo
weak armor	Raises Speed and lowers Defense by one stage each upon being hit by a physical move.	kabutops
water veil	Prevents burns.	wailord
water veil	Prevents burns.	seaking
water veil	Prevents burns.	mantine
water veil	Prevents burns.	lumineon

As you can see in the picture, the query returns 595 rows (too many to fit in a screenshot).

5.3.2 Explanation

In here we were asked to provide the *name* of the ability, its corresponding *description* and the *species name* of those species whose starting *base experience* is greater than the average *base experience* of all the species. To obtain this, we need to *JOIN species*, *has_ability* and *ability* tables. In the *WHERE* clause we specify that we want to display the information of only those species whose starting *base_experience* is greater than the

average of the *basic_experience* of all the species. Since the AVG function gets the *basic_experience* of all the species, a subquery is needed.

5.3.3 Query validation

In order to check that the query returns what is expected, we can get the average of the *basic_experience* of all the species. Once we have that, we can create another query that shows the species with a basic experience greater than this average value. If we start checking name by name we see that it matches with our initial query.

```
SELECT AVG(basic_experience) AS average
FROM species;
```

average
145.9553752535496957

```
SELECT s.name, s.basic_experience
FROM species AS s
WHERE s.basic_experience > 145.95
ORDER BY s.basic_experience DESC;
```

#	name	basic_experience
1	blissey	635
2	chansey	395
3	rayquaza	340
4	ho-oh	340
5	giratina-altered	340
6	palkia	340
7	dialga	340
8	lugia	340
9	mewtwo	340
10	kyogre	335

5.4 Query 4

5.4.1 Solution

```
SELECT stat_name, max_base_stat, evolved_species
FROM (
    SELECT evolved_species.name AS evolved_species, bs.name AS stat_name,
    bs.base_stat AS max_base_stat,
        ROW_NUMBER() OVER (PARTITION BY bs.name ORDER BY bs.base_stat DESC,
    evolved_species.id) AS rn
    FROM Species AS baby_species
    JOIN evolve ON baby_species.id = evolve.id_current_species
    JOIN Species AS evolved_species ON evolve.id_next_species = evolved_species.id
    JOIN Pokemon_Creature ON baby_species.id = Pokemon_Creature.species
    JOIN has_stat ON evolved_species.id = has_stat.id_species
    JOIN Base_Statistics AS bs ON has_stat.id_statistics = bs.id
    WHERE Pokemon_Creature.is_baby = TRUE
) AS subquery
WHERE rn = 1
```

```
ORDER BY max_base_stat DESC, stat_name;
```

	stat_name	max_base_stat	evolved_species
1	hp	250	chansey
2	special defense	140	mantine
3	attack	120	hitmonlee
4	defense	115	sudowoodo
5	special attack	115	jynx
6	speed	105	electabuzz

5.4.2 Explanation

This query obtains the maximum base stat of each statistics among the evolved species of baby pokemon. It uses a subquery to *JOIN* first the tables **species** with **evolve** and again with **species** in order to know the evolution from a baby, and then, with **pokemon_creature**, **has_stat** and **base_statistics** to know the *maximum base stat*. Note that we need the **pokemon_creature** table because we stored if it is a baby or not there.

The subquery also uses row number () that calculates for each combination of statistics and evolved species. This number, which we call rn, is determined by ordering the **basic_statistics** in descending order and breaking ties according to the evolved species ID.

Then, PARTITION BY is used to ensure that the row number is calculated for each **base_statistics**.

The main query filters the subquery results to only include rows with rn = 1(meaning the row with the highest **base_stat** value for each **base_statistics**). Finally, the output is ordered by **base_stat** value in descending order.

5.4.3 Query validation

If we check all the max base statististics of each of the pokemons, we can see if some are repeated and should appear twice.

evolved_species	stat_name	max_base_stat	rn
hitmonlee	attack	120	1
hitmonlee	attack	120	34
snorlax	attack	110	35
sudowoodo	defense	115	1
sudowoodo	defense	115	30
hitmontop	defense	95	31
1105 chansey	hp	250	1
chansey	hp	250	22
wobuffet	hp	190	23

jynx	special attack	115	1
jynx	special attack	115	28
lucario	special attack	115	29
mantine	special defense	140	1
mantine	special defense	140	25
mr-mime	special defense	120	26
electabuzz	speed	105	1
electabuzz	speed	105	27
jynx	speed	95	28

With this results we can see that the max base statistic of each of them is correct and that the only ones that have the same max stat are jynx and lucario (special attack).

5.5 Query 5

5.5.1 Solution

```
(SELECT
    'Strength' AS type, attacker.name AS strongest_type, COUNT(*) AS
strength_count
FROM affect AS a
JOIN Type AS attacker ON attacker.name = a.name_type1
WHERE a.multiplier = 'x2'
GROUP BY attacker.name
ORDER BY strength_count DESC
LIMIT 2)

UNION ALL

(SELECT
    'Resistance' AS type, defender.name AS resistant_type, COUNT(*) AS
resistance_count
FROM affect AS a
JOIN Type AS defender ON defender.name = a.name_type2
WHERE a.multiplier = 'x0.5'
GROUP BY defender.name
ORDER BY resistance_count DESC
LIMIT 1);
```

	type	strongest_type	strength_count
1	Strength	ground	5
2	Strength	fighting	5
3	Resistance	steel	10

5.5.2 Explanation

This query provides us information about the strongest attacking type and the most resistant defending type. In the first part of the query, we find the attacking type (strongest type) based on the count of moves (strength_count) that have a multiplier of x2 in the **affect** table. It groups the results by attacker name and orders them in descending order of strength_count. It is limited to 2 since there are two types that have the maximum strength.

In the second part of the query, it identifies the most resistant defending type (resistant type) based on the count of moves (resistance_count) that have a multiplier of x0.5 in the **affect** table. It groups the results by defender name and orders them in descending order of resistance_count. It is just limited to 1 since there is only one type with the most resistance.

A *UNION ALL* is used since we want to combine the results of both queries. It basically adds a new column called ‘type’ that indicates whether the row represents the strongest type or the most resistant type.

5.5.3 Query validation

The validation of this query can be done by running the queries separately.

	strongest_type	strength_count		resistant_type	resistance_count
1	ground	5		steel	10
2	fighting	5		fire	6
3	ice	4		poison	5
4	fire	4		dragon	4
5	rock	4		grass	4
6	fairy	3		water	4
7	water	3		rock	4
8	flying	3		fairy	3
9	grass	3		flying	3
10	steel	3		fighting	3
11	bug	3		electric	3
12	ghost	2		bug	3
13	electric	2		ghost	2
14	dark	2		dark	2
15	psychic	2		ground	2
16	poison	2		psychic	2
17	dragon	1		ice	1

As we can see, there are two strength types with the maximum number of strengths and just one with the maximum resistances.

5.6 Query 6

5.6.1 Solution

```
SELECT scl.id AS species_id, scl.name AS species_name, scl.basic_experience AS species_experience,
       sc2.id AS evolved_species_id, sc2.name AS evolved_species_name,
       (scl.basic_experience + sc2.basic_experience) AS evolved_species_experience,
       sc3.id AS final_evolved_species_id, sc3.name AS final_evolved_species_name,
       (scl.basic_experience + sc2.basic_experience + sc3.basic_experience) AS final_evolved_species_experience
FROM Species AS scl
```

```

JOIN evolve AS ev1 ON scl.id = ev1.id_current_species
JOIN Species AS sc2 ON ev1.id_next_species = sc2.id
JOIN evolve AS ev2 ON sc2.id = ev2.id_current_species
JOIN Species AS sc3 ON ev2.id_next_species = sc3.id
WHERE ev1.trigger = 'level up'
AND ev1.level <> 0
AND ev2.trigger = 'level up'
AND ev2.level <> 0
ORDER BY final_evolved_species_experience DESC
LIMIT 1;

```

	species_id	species_name	species_experience	evolved_species_id	evolved_species_name	evolved_species_experience
1	147	dratini	60	148	dragonair	207
				final_evolved_species_id	final_evolved_species_name	final_evolved_species_experience
				149	dragonite	507

5.6.2 Explanation

Now we want to get the evolution chain that requires more experience to reach its biggest form. To reach this goal, we select the id of the species, the species name and the species experience. We do this also for the evolved and for the final evolved species. We **JOIN** **species** with **evolve** multiple times since we want to know the evolution relationships. We link the *species.id* with the *current_species*, the *next_species* with the *second species id*, the *second species id* with the *current species of the second evolution* and finally, the *next_species id of the second evolution* with the *third species id*.

In the WHERE clause, we state that in order to evolve, the trigger in evolution must be 'level up' and that the level is different from 0.

5.6.3 Query validation

In this case, we have UPDATED the *basic_experience* of bulbasaur to 1 million in order to see it we get that as the biggest evolution chain.

```

UPDATE species
SET basic_experience = 10000000
WHERE id = 1;

```

After the modification, we see that we get bulbasur.

	species_id	species_name	species_experience	evolved_species_id	evolved_species_name	evolved_species_experience
1	1	bulbasaur	10000000	2	ivysaur	10000142
				final_evolved_species_id	final_evolved_species_name	final_evolved_species_experience
				3	venusaur	10000405

5.7 Trigger 1

5.7.1 Solution

```

DROP TABLE IF EXISTS warnings CASCADE;
CREATE TABLE warnings (
    affected_table VARCHAR(30) NOT NULL,
    date DATE NOT NULL,
    error_message VARCHAR(255) NOT NULL,

```

```

    "user" VARCHAR(255) NOT NULL
);

-- Trigger maximum num types x species
DROP FUNCTION IF EXISTS check_max_types CASCADE;
CREATE OR REPLACE FUNCTION check_max_types() RETURNS TRIGGER AS $$

DECLARE
    types_count INT;
    total_entries INT;
BEGIN
    SELECT COUNT(id_species) INTO types_count
    FROM has_type
    WHERE id_species = NEW.id_species;

    total_entries := types_count + 1;

    IF types_count > 2 THEN
        INSERT INTO warnings (affected_table, date, error_message, "user")
            VALUES ('has_type', CURRENT_DATE, 'A ' || (total_entries - 1) || 'th
entry has been inserted into the has_type table', current_user);
    END IF;

    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

DROP TRIGGER IF EXISTS max_types_trigger ON has_type CASCADE;
CREATE TRIGGER max_types_trigger
AFTER INSERT OR UPDATE ON has_type
FOR EACH ROW
EXECUTE FUNCTION check_max_types();

-- Trigger maximum num abilities x species
DROP FUNCTION IF EXISTS check_max_abilities CASCADE;
CREATE OR REPLACE FUNCTION check_max_abilities() RETURNS TRIGGER AS $$

DECLARE
    abilities_count INT;
    total_entries INT;
BEGIN
    SELECT COUNT(id_ability) INTO abilities_count
    FROM has_ability
    WHERE id_species = NEW.id_species AND hidden = FALSE;

    total_entries := abilities_count + 1;

    IF abilities_count > 2 THEN
        INSERT INTO warnings (affected_table, date, error_message, "user")
            VALUES ('has_ability', CURRENT_DATE, 'A ' || (total_entries - 1) || 'th
entry has been inserted into the has_ability table', current_user);
    END IF;

    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

DROP TRIGGER IF EXISTS max_abilities_trigger ON has_ability CASCADE;
CREATE TRIGGER max_abilities_trigger
AFTER INSERT OR UPDATE ON has_ability
FOR EACH ROW
EXECUTE FUNCTION check_max_abilities();

```

5.7.2 Explanation

This triggers are basically activated after inserting or updating the tables **has_ability** and **has_type**. Two variables are declared, one that counts the number of abilities or types that a specific species has and the other one to notice the number of entries that have been inserted into the table. If the number of either types or abilities per species is greater than 2, an INSERT INTO the table warnings is done warning that x entries have been inserted into the **has_ability** table. In this warnings table, the affected table, the current date of the insertion and the user who did that operation is also registered.

5.7.3 Query validation

Since there are two different triggers, we must do two different validations, one for the types and one for the abilities.

For the first trigger to activate, we first check how many types a particular species has. In this case, we are going to use the species id 493. We go to the **has_type** table and see that it has only one type associated with it. Since the trigger is only triggered when it detects more than two different types associated to the same species, the next step is to insert two more types into the **has_type** and **type** tables.

```
INSERT INTO Type (id, name)
VALUES (22, 'ball');
INSERT INTO has_type (id_species, type_name)
VALUES (1, 'ball');
```

The figure consists of three vertically stacked screenshots of a PostgreSQL terminal window. The top screenshot shows the creation of a new type 'ball' with ID 22. The middle screenshot shows the 'has_type' table for species ID 493, which contains one entry ('ball') and is not yet triggering the trigger. The bottom screenshot shows the 'has_type' table after two more entries ('force' and 'normal') have been added, resulting in three entries total, which triggers the trigger.

id	type_name
1	ball
2	force
3	normal

After each insertion we run the trigger and check that in the second insertion nothing happens while in the third one, the corresponding error message appears.

A screenshot of a PostgreSQL terminal showing a single row in a table named 'affected_table'. The table has columns for 'affected_table', 'date', 'error_message', and 'user'. The row shows that a 3rd entry was inserted into the 'has_type' table by the 'postgres' user on May 28, 2023.

affected_table	date	error_message	user
has_type	2023-05-28	A 3th entry has been inserted into the has_type table	postgres

We repeat the same process but this time inserting new abilities in the table **has_ability** and **ability**.

5.8 Trigger 2

5.8.1 Solution

5.8.2 Explanation

5.8.3 Query validation

We faced some problems when doing this trigger and didn't have enough time to fix it.

6 I'm not a gamer, I'm a Pokémon trainer

6.1 Query 1

6.1.1 Solution

```
SELECT tr.name AS trainer, pc.id AS firstPokemon
    FROM trainer tr JOIN team t ON t.trainer = tr.id
    JOIN battle b ON t.id = b.winner JOIN team glt ON glt.id = b.loser
    JOIN gym_leader gl ON glt.trainer = gl.id
    JOIN pokemon_creature pc ON pc.trainer = tr.id
    GROUP BY tr.id, pc.id
    ORDER BY COUNT(gl.id) DESC, MIN(pc.date_time)
    LIMIT 1;
```

	trainer	firstpokemon
1	Olinda	6167

6.1.2 Explanation

To find the trainer who had defeated the most Gym Leaders in combat we had to select its name from the table **Trainer**, then join this one with its **Team** to see in which battles he had fought. We were able to get only the battles in which they won by linking its team to the *winner* attribute in the **Battle** entity. To find only those battles in which the trainer had defeated a Gym Leader, we joined the *loser* attribute of the **Battle** entity with a new instance of the **Team** entity. Then, we linked this team to the **Gym_Leader entity** (we could skip the **Trainer** table as the **Gym_Leader** has as its *id* a FK to the **Trainer** table).

The above mentioned steps return us all the trainers that have defeated a Gym Leader. To get the amount of wins of each trainer, we used a *GROUP BY* clause to join all the rows from the same trainer. To order them by the one that had defeated most, we used an *ORDER BY* clause in which we specified we wanted the output to be from the one with a highest count of defeated Gym Leaders (*COUNT* operation of Gym Leaders) to the one that had the least (*DESC*). We added a *LIMIT 1* to only get the trainer that had defeated the most.

Lastly, to show the first Pokémon's id of the trainer, we joined the **Pokemon_Creature** entity with the **Trainer** (adding it on the *GROUP BY* clause) and ordering the output using the *ORDER BY* clause (as a second parameter, giving it less priority) with the Pokémon that had been captured earlier. To get this, we used the *M/N* operation to check which of their Pokémons had the smallest *date_time* attribute.

6.1.3 Query validation

It is hard to check that this query works with the given dataset, as if we remove the *LIMIT* of the previous query and show the *COUNT* of defeated Gym Leaders, and the *date_time* of when the Pokémon was captured, we can see how all trainers have defeated a maximum of one Gym Leader. Olinda (output of the query) just happens to have the Pokémon who was captured the earliest of all (second condition of the *ORDER BY* clause). Therefore, she is returned as the output.

```
SELECT tr.name AS trainer, pc.id AS firstPokemon, pc.date_time AS captured,
COUNT( gl.id) AS defeated
```

```

FROM trainer tr JOIN team t ON t.trainer = tr.id
JOIN battle b ON t.id = b.winner JOIN team glt ON glt.id = b.loser
JOIN gym_leader gl ON glt.trainer = gl.id
JOIN pokemon_creature pc ON pc.trainer = tr.id
GROUP BY tr.id, pc.id ORDER BY COUNT(gl.id) DESC, MIN(pc.date_time);

```

	trainer	firstpokemon	captured	defeated
1	Olinda	6167	1996-08-18 19:17:20.000000	1
2	Garrett	12469	1996-09-08 10:33:00.000000	1
3	Karol	829	1996-09-14 07:53:04.000000	1
4	Falkner	10490	1996-09-28 12:27:59.000000	1
5	Nannie	6020	1996-10-08 09:26:24.000000	1
6	Rozanne	6375	1996-12-18 11:06:06.000000	1
7	Garrett	12460	1996-12-21 04:25:18.000000	1
8	Mac	6865	1996-12-26 10:30:57.000000	1

To check if our query actually works, we can try manually inserting some data. Let's see we want the trainer *Karol* to defeat another Gym Leader. We will insert the following values to the **Battle** table:

```

INSERT INTO battle(id, winner, loser)
VALUES (700, 177, 424);

```

As you can see on the screenshots below, the *id* 177 corresponds to the trainer *Karol* (*id* = 30) and we set this one as the winner of the battle. Then, we can also identify the team 428 as property of the trainer 504 (first gym leader) and we set it as the loser.

Karol's team: (**Trainer** and **Team** tables respectively)

id	class	name	experience	gold
30	Spy	Karol	121717	-109

id	trainer
177	30

Gym Leader's team: (**Gym_leader** and **Team** tables respectively)

id	gym
504	527

id	trainer
424	504

Therefore, if we run the query again we get the following output:

	trainer	firstpokemon
1	Karol	829

As you can see, now the trainer that appears is Karol, as these one has defeated 2 Gym Leaders in total.

6.2 Query 2

6.2.1 Solution

```
SELECT pc.nickname AS pokemon, tr.name AS trainer, COUNT(b.id) AS defeats
    FROM pokemon_creature pc JOIN fight f ON f.id_creature = pc.id
    JOIN battle b ON f.id_battle = b.id
    JOIN trainer tr ON tr.id = pc.trainer
    WHERE b.winner = pc.team AND f.life_remaining = 0
    GROUP BY pc.id, tr.id
    ORDER BY COUNT(b.id) DESC, MIN(pc.date_time)
    LIMIT 5;
```

	pokemon	trainer	defeats
1	Mac	Norberto	5
2	Rhett	Wiley	4
3	Friday	Mark	4
4	Nibbles	Wiley	4
5	Ginger	Karoline	4

6.2.2 Explanation

To get the Pokémons that have been defeated the most times in battles they have won we have started by joining the tables **Pokemon_Creature** and **Battle** through the N:M relation table **Fight** (which holds the statistics of the Pokémons after each battle). With this, we get all the Pokémons and the battles they have fought in. To identify the ones they won we add a *WHERE* clause linking the Pokémons's *team* to the *winner* attribute of the **Battle** entity.

Also, to get the battles in which they were particularly defeated (even though their team won) we add a second condition to the *WHERE* clause stating that their hp after finishing the battle (*life_reamining* attribute in **Fight** table) has to be equal to 0.

To show the ones that have been defeated the most times, we add a *GROUP BY* clause to join rows with the same Pokémons and trainer (to count them) and add a *GROUP BY* condition to order the output from the one with most defeats (using the *COUNT* operator to count the number of battles) to the one that has less (DESC). We also limit the query to only show five Pokémons using the *LIMIT* operator.

Lastly, to give priority to the Pokémons that were captured the earliest, we add a second condition to the *ORDER BY* clause (less priority) that indicates we want to order them by the minimum capture time (using the *MIN* operator and accessing the *date_time* attribute of the **Pokemon_Creature** entity). We also have to do a *JOIN* with the **Trainer** entity to display the Pokémons trainer's name.

6.2.3 Query validation

First, we will check if the results we got match the criteria by displaying the *HP* attribute of the Pokémons, their team *id* (so that we can later check if the count of lost battles makes sense), and the *date_time* when the Pokémons was captured. By executing the following query, we get the result below.

```
SELECT pc.id, pc.nickname AS pokemon, tr.name AS trainer, COUNT(b.id) AS defeats, pc.hp AS hp, pc.team AS team, pc.date_time AS captured
```

```

    FROM pokemon_creature pc JOIN fight f ON f.id_creature = pc.id
    JOIN battle b ON f.id_battle = b.id
    JOIN trainer tr ON tr.id = pc.trainer
    WHERE b.winner = pc.team AND f.life_remaining = 0
    GROUP BY pc.id, tr.id ORDER BY COUNT(b.id) DESC, MIN(pc.date_time) LIMIT
5;

```

	id	pokemon	trainer	defeats	hp	team	captured
	3575	Mac	Norberto	5	0	189	2005-03-13 04:16:39.000000
	4940	Rhett	Wiley	4	0	436	2000-02-17 13:00:11.000000
	425	Friday	Mark	4	0	115	2004-11-14 14:00:16.000000
	4939	Nibbles	Wiley	4	0	436	2007-01-03 13:10:57.000000
	8111	Ginger	Karoline	4	0	391	2007-08-22 16:35:04.000000

As you can see, all the Pokémons have an *HP* of zero, matching the criteria specified by the statement. Moreover, see that for those Pokémons that have been defeated the same number of times (4), these ones are ordered by the ones that were captured before (starting in 2000 and going up to 2007).

To validate the count of battles is correct, we will take the first Pokémon as an example (*id* of 3575 and *team* with *id* 189), and execute the following query which will display all the battles in which this Pokémon's team won but he had 0 HP:

```

SELECT pc.nickname, b.id
    FROM pokemon_creature pc JOIN fight f ON pc.id = f.id_creature
    JOIN battle b ON f.id_battle = b.id
    WHERE f.life_remaining = 0 AND pc.team = b.winner AND pc.team = 189
    AND pc.id = 3575;

```

	nickname	id
1	Mac	326
2	Mac	180
3	Mac	470
4	Mac	469
5	Mac	457

As you can see in the screenshot above, this query retrieved five rows, which is exactly the count of battles displayed on the previous query. Meaning that, our count of battles works correctly and is able to identify the desired Pokémon.

6.3 Query 3

6.3.1 Solution

```

SELECT tr.id AS trainer, COUNT(pc.id) AS numpokemon
    FROM trainer tr JOIN team t ON tr.id = t.trainer
    LEFT JOIN battle b ON t.id = b.winner
    JOIN pokemon_creature pc ON pc.trainer = tr.id
    WHERE b.winner IS NULL AND tr.experience > tr.gold^2

```

```
GROUP BY tr.id;
```

	trainer	numpokemon
1	426	10
2	59	26
3	65	27
4	121	41
5	239	38
6	74	33
7	34	12
8	67	43
9	484	17
10	36	21
11	277	31
12	128	33
13	432	40

6.3.2 Explanation

To show the trainers that had never won a battle we did a *LEFT JOIN* with the **Battle** table (connecting the *id* of the **Trainer** with the *winner* attribute of the **Battle** entity) and specified in the *WHERE* clause that the *winner* attribute had to be *NULL*. We used a *LEFT JOIN* for this so that we could get the non-existing relationships between these two tables.

Also, to narrow down even more our output and display only those trainers that have more experience than squared amount of gold we added a second *WHERE* clause in which we specified the relationship between these two attributes.

Lastly, to count the number of Pokémons each trainer owns, we did a *JOIN* with the **Pokemon_Creature** entity (linking the *trainer* attribute of this one with the **Trainer's id**) and added a *GROUP BY* clause to join all the rows of the same trainer. This way, we could use the *COUNT()* operator to check how many Pokémons the trainer had.

6.3.3 Query validation

If execute the same query as before, but now displaying the *experience* and *gold* attributes from the **Trainer** entity, we will be able to check if the relationship between these is the expected:

```
SELECT tr.id AS trainer, COUNT(pc.id) AS numpokemon, tr.experience, tr.gold
  FROM trainer tr JOIN team t ON tr.id = t.trainer
  LEFT JOIN battle b ON t.id = b.winner
  JOIN pokemon_creature pc ON pc.trainer = tr.id
 WHERE b.winner IS NULL AND tr.experience > tr.gold^2
 GROUP BY tr.id;
```

	trainer	numpokemon	experience	gold
1	426	10	365665	308
2	59	26	537855	-275
3	65	27	494951	-557
4	121	41	1880445	351
5	239	38	348903	192
6	74	33	1821347	-563
7	34	12	1888246	1283
8	67	43	1672810	882
9	484	17	1451825	-342
10	36	21	1875336	885
11	277	31	1347591	378
12	128	33	1237790	419
13	432	40	329603	387

As you can see, all the trainers displayed have more experience than squared their amount of gold.

Now, if we insert a new Pokémon to one of the trainers displayed (the first one for example) we can see how the count of Pokémons changes.

```
INSERT INTO pokemon_creature(id, trainer) VALUES (15000, 426);
```

	trainer	numpokemon	experience	gold
1	426	11	365665	308
2	59	26	537855	-275
3	65	27	494951	-557

You can see how the count was increased by one unit on the first trainer. Meaning that we were able to correctly count the number of Pokémons each trainer has.

To ensure the displayed trainers match the criteria of never having won a battle, we will take the first displayed trainer (*id* = 426) and show all the battles they have been in, and see if they won or lost:

```
SELECT tr.id AS trainer, bw.winner, bl.loser AS loser
  FROM trainer tr JOIN team t ON tr.id = t.trainer
  LEFT JOIN battle bl ON t.id = bl.loser
  LEFT JOIN battle bw ON t.id = bw.winner
 WHERE tr.id = 426;
```

	trainer	winner	loser
1	426	<null>	283
2	426	<null>	283
3	426	<null>	283
4	426	<null>	283

As we can see in the screenshot above, the trainer with $id = 426$ has fought in four battles, but he didn't win any (*winner* attribute is null). Notice how the *loser* attribute matches the trainer's team:

id	trainer
283	426

Therefore, we can conclude that the displayed trainer matched the criteria and are giving us the expected result.

6.4 Query 4

6.4.1 Solution

```
SELECT o.name AS organization, t.name AS villain, ec.money_stole AS stolen
  FROM criminal_organization o
  JOIN evil_criminal ec ON o.name = ec.criminal_organization
  JOIN trainer t ON t.id = ec.id
 WHERE ec.money_stole IN
  (SELECT evc.money_stole
    FROM evil_criminal evc
   WHERE o.name = evc.criminal_organization
   ORDER BY evc.money_stole DESC
   LIMIT 2)
  ORDER BY o.name, money_stole DESC;
```

	o.name	t.name	money_stole
1	Aqua	Mia	1581
2	Aqua	Willian	1576
3	Cipher	Carlton	1497
4	Cipher	Jaye	1398
5	Galactic	Jade	1669
6	Galactic	Erin	1535
7	Magma	Bennett	1530
8	Magma	Kerri	1399
9	Plasma	Ghetis	159
10	Rocket	Giovanni	6969700000000000
11	Rocket	Fidel	1318

6.4.2 Explanation

To show the villains from each villainous organization we have joined the tables **Criminal_Organization** and **Evil_Criminal** (linking the *criminal_organization* attribute of the criminal with the organization). To get their names, we have also joined the **Trainer** entity (equaling the *id* of the evil criminal with the trainer's).

Moreover, to get the two criminals that have stolen the most money we had to do a subquery that returned the two criminals from each organization with the most stolen money and use it in the *WHERE* clause. We had to take this approach as there was no other way for us to check if the stolen money was one of the desired values. If the query

had only wanted the criminal with the most stolen money, we could have omitted this subquery and just use a *LIMIT 1* operator.

This subquery selected the stolen money of the two evil criminals that belonged to the organization being evaluated (specified in the inner *WHERE* clause), ordered them from the greatest to the lowest (*DESC*) and limited the output to only 2 (using the *LIMIT* operator).

To show the output with the desired format, we used an *ORDER BY* expression to force the output to show the names of the criminal organizations in alphabetical order and (as a secondary condition) from the greatest money stolen to the least (*DESC*).

6.4.3 Query validation

To validate this query, we will display all the criminals from a criminal organization and sort them by the one with the most stolen money to the least using the following query:

```
SELECT o.name, t.name, ec.money_stole
      FROM criminal_organization o
      JOIN evil_criminal ec ON o.name = ec.criminal_organization
      JOIN trainer t ON t.id = ec.id
 ORDER BY ec.money_stole DESC;
```

	o.name	t.name	money_stole
1	Rocket	Giovanni	696970000000000
2	Galactic	Jade	1669
3	Aqua	Mia	1581
4	Aqua	Willian	1576
5	Aqua	Gregorio	1567
6	Galactic	Erin	1535
7	Magma	Bennett	1530
8	Aqua	Lela	1502
9	Cipher	Carlton	1497
10	Aqua	Ralph	1472
11	Galactic	Marisol	1458
12	Aqua	Hermine	1417
13	Magma	Kerri	1399
14	Cipher	Jaye	1398
15	Magma	Preston	1350
16	Cipher	Jonah	1326
17	Rocket	Fidel	1318
18	Aqua	Douglass	1308

87 Plasma Ghetsis 159

As you can see, the output matches with the one on the previous section. Notice how the Plasma criminal organization only gave us one row on the previous output. We will check how many evil criminals this one has with the following query:

```
SELECT ec.id FROM criminal_organization co
      JOIN evil_criminal ec ON co.name = ec.criminal_organization
     WHERE co.name = 'Plasma';
```

<	<	1 row	>	>
	id			
1	535			

As you can see, there is only one trainer in the criminal organization called ‘Plasma’. This explains why we only got one row for this organization in the previous outputs.

Now we will try to insert a new evil criminal into the ‘Aqua’ organization and give it a greater number of money stolen with the following query:

```
INSERT INTO evil_criminal(id, money_stole, criminal_organization)
VALUES (1, 9999999, 'Aqua');
```

We execute the initial query again and we get the following:

	organization	villain	stolen
1	Aqua	Genie	9999999
2	Aqua	Mia	1581
3	Cipher	Carlton	1497
4	Cipher	Jaye	1398

As you can see, the villains being displayed for the ‘Aqua’ organization changed and now include this new criminal, as it has a greater number of stolen money.

6.5 Query 5

6.5.1 Solution

```
SELECT pc.nickname AS pokemon, e.nickname AS enemy, pc.status_condition AS ailment
FROM pokemon_creature pc JOIN fight f ON f.id_creature = pc.id
JOIN battle b ON f.id_battle = b.id JOIN fight fe ON fe.id_battle = b.id
JOIN pokemon_creature e ON fe.id_creature = e.id
JOIN cause_or_receive cor ON e.id = cor.creature
JOIN status s ON cor.move = s.id
WHERE pc.status_condition NOT LIKE 'none' AND e.id != pc.id
AND s.special_effect = pc.status_condition AND pc.team != e.team
GROUP BY pc.id, e.id, s.special_effect;
```

	pokemon	enemy	ailment
1	Maya	Snowy	confusion
2	Jackie	Cojoncio	sleep
3	Salty	Rambo	confusion
4	Tramp	Tiny	poison
5	Buffie	Tiny	poison
6	Fuzzy	Hamlet	freeze
7	Fuzzy	Kramer	freeze
8	Bunky	Thelma	poison
9	Wilson	Aries	confusion
10	Wilson	Porter	confusion
11	Cupcake	Annie	sleep
12	Bridgett	Boo	paralysis
13	Bridgett	Holly	paralysis
14	Julius	Ben	paralysis
15	Julius	Flint	paralysis
16	Mini	Bubbles	confusion

As you can see in the picture, the query returns 399 rows (too many to fit in the screenshot).

6.5.2 Explanation

Firstly, to show all the Pokémons that are under an ailment state we selected the *nickname* of the **Pokemon_Creature** table and added in a *WHERE* clause that its *status_condition* attribute could not be ‘none’. Next, to get all the battles they fought in we joined the **Pokemon_Creature** and **Battle** tables (by linking them with the **Fight** entity).

Now, to get all the Pokémons they fought against we did the same joins as before but in an inverse order (linking the **Battle id** of the ones gotten before to a new **Fight** instance and a new **Pokemon_Creature** to get the Pokémons that also fought on those battles). To ensure we don’t get duplicates (same Pokémon as it did fight in the battle) or Pokémons from the same team as the one we are evaluating, we added two conditions in the *WHERE* clause specifying that the two Pokémons had to be different (*pc.id != e.id*) and that they couldn’t belong to the same team (*pc.team != e.team*).

Lastly, instead of getting all the Pokémons they fought against, we join these Pokémon with the **Status** (move) table (through the **Cause_Or_Receive** table). This way, we can add an extra condition in the *WHERE* clause specifying that the Pokémon must inflict an ailment of the same status as the affected Pokémon is in, discarding like this all those Pokémon that couldn’t have caused the damage.

To display the information, we add a *GROUP BY* clause that joins those rows that have the same two Pokémons with the same special effect (eliminating repetitions for those Pokemons that have two moves of the same type).

6.5.3 Query validation

To validate this query we will take the first displayed Pokémon of an example (Maya). We will display the query again showing the Pokémon *id* so we can use in further queries and only Maya's Pokemons by adding in the *WHERE* clause that the Pokémons *id* has to be 134:

	<input type="checkbox"/> id	<input type="checkbox"/> pokemon	<input type="checkbox"/> enemy	<input type="checkbox"/> ailment
1	134	Maya	Snowy	confusion

Now that we know Maya's *id* and have seen that there is only one row retrieved by the query, we will execute the following query to show all the Pokémon Maya has encountered in all its battles:

```
SELECT pc.id, pc.nickname, e.nickname, e.id, s.special_effect
  FROM pokemon_creature pc JOIN fight f ON f.id_creature = pc.id
    JOIN battle b ON f.id_battle = b.id JOIN fight fe ON fe.id_battle = b.id
      JOIN pokemon_creature e ON fe.id_creature = e.id
      JOIN cause_or_receive cor ON e.id = cor.creature
      JOIN status s ON cor.move = s.id
 WHERE pc.id = 134
 ORDER BY s.special_effect;
```

	<input type="checkbox"/> pc.id	<input type="checkbox"/> pc.nickname	<input type="checkbox"/> e.nickname	<input type="checkbox"/> e.id	<input type="checkbox"/> special_effect
1	134	Maya	Lexus	12328	burn
2	134	Maya	Maximus	8128	burn
3	134	Maya	Winter	8842	burn
4	134	Maya	Winter	8842	burn
5	134	Maya	Winter	8842	burn
6	134	Maya	Lexus	12328	burn
7	134	Maya	Snowy	8843	confusion
8	134	Maya	Bumper	9320	freeze
9	134	Maya	Atlas	8130	freeze
10	134	Maya	Karma	132	infatuation

As you can see, there is only one Pokémon (Snowy) that has a special effect that causes *confusion*, which we saw was the ailment state that Maya is in. Therefore, we can conclude that the output retrieved by the initial query makes sense as it only retrieved the one out of 40 Pokémons encountered that have a move that inflicts their ailment state.

Just to double check that the *Snow* Pokémon has that specific move, we will display all the moves that this Pokémon (*id* = 8843) has with the following query:

```
SELECT m.name, s.special_effect
  FROM move m JOIN cause_or_receive cor ON m.id = cor.move
    LEFT JOIN status s ON m.id = s.id
    JOIN pokemon_creature pc ON cor.creature = pc.id
 WHERE pc.id = 8843;
```

	name	special_effect
1	karate chop	<null>
2	rolling kick	<null>
3	triple kick	<null>
4	dynamic punch	confusion

As you can see, this PokéMon has four moves, out of which one can inflict an ailment state, which is in fact *confusion*, the ailment the Maya PokéMon has.

6.6 Query 6

6.6.1 Solution

```
SELECT ty.name, COUNT(b.loser)
    FROM team t LEFT JOIN battle b ON b.loser = t.id
    JOIN gym_leader gl ON gl.id = t.trainer JOIN gym g ON gl.gym = g.id
    JOIN type ty ON g.id_type = ty.id
    GROUP BY ty.name ORDER BY COUNT(b.loser)
    LIMIT 1;
```

	name	count
1	dragon	0

6.6.2 Explanation

To get the Gym Leaders that belong to each Gym type we first join the **Gym_Leader** and **Gym** entities to get the gym each leader has. Then, we join the **Gym** and **Type** tables to get the type of the leader's gyms. Lastly, we use the **GROUP BY** clause to join all those trainers that belong to the same gym type. With this, we are able to use the **COUNT** operator to check how many Gym Leaders belong to each gym type.

To figure out which of the selected types is the hardest to beat, we check which has the least number of defeated trainers. To do so, we first link the **Team** and **Gym_Leader** entities to then be able to join the **Team** table with the *loser* attribute of the **Battle** entity (we use a *LEFT JOIN* for this relationship as we want to allow *NULL* relations to count those types that don't have any leader who has been defeated). Now we can use the **COUNT** operator to check how many losers each type has. We also use this operation in an **ORDER BY** clause so that we can sort them by the ones with the least number of defeats. Lastly, we add a *LIMIT 1* to only get the type with the least number of defeats.

6.6.3 Query validation

As happened in other triggers, there are many rows that satisfy the given conditions, as you can see with the following query. If we use the same query but take out the *LIMIT* we get the following:

	name	count
1	dragon	0
2	ice	0
3	grass	0
4	psychic	0
5	flying	1
6	rock	1
7	fire	1
8	normal	1
9	ground	1
10	water	1
11	fighting	1
12	steel	2
13	electric	2
14	ghost	3
15	bug	3
16	poison	4

As you can see, there are four types whose leaders have not been defeated even once. Therefore, to check that our query actually works we will insert three new battles which define trainers from the ‘dragon’, ‘ice’ and ‘grass’ types as losers.

First, we identify these type’s *id* on the **Type** table and find a gym that identifies with that type in the **Gym** table:

	id	name
	12	grass
	15	ice
	16	dragon

	id	medal	id_item	id_type
	553	Forest Badge	306	12
	551	Glacier Badge	352	15
	545	Rising Badge	329	16

We will now find the gym leaders that own these three gyms.

	id	gym
	525	553
	514	551
	515	545

Finally, let’s find these trainer’s teams in the **Team** table:

	id	trainer
	321	525
	413	514
	468	515

Now we execute the following query to insert three new battles defining the found teams as losers and the team with *id* = 196 as the winner.

```
INSERT INTO battle(id, winner, loser) VALUES (650, 196, 321);
INSERT INTO battle(id, winner, loser) VALUES (651, 196, 413);
INSERT INTO battle(id, winner, loser) VALUES (652, 196, 468);
```

name	count
psychic	0
normal	1
ground	1
water	1
flying	1
dragon	1
fire	1
ice	1

As you can see, if we run the query again we now get that the count of defeats has increased for the selected types, having 'psychic' as the only type without any defeated gym leader. This proves that the counting of the defeats works and that the query would be able to identify the type with less defeated leaders.

6.7 Trigger 1

6.7.1 Solution

```
DROP FUNCTION IF EXISTS teach_move CASCADE;
CREATE OR REPLACE FUNCTION teach_move() RETURNS TRIGGER AS $$
BEGIN
    IF (
        EXISTS
            (SELECT *
                FROM trainer t JOIN pokemon_creature pc ON pc.trainer = t.id
                JOIN owns_items oi ON t.id = oi.trainer_id
                JOIN move_machine mm ON oi.item_id = mm.id
                WHERE mm.move = NEW.move AND pc.id = NEW.creature
            )
    ) THEN
        --remove move machine
        DELETE FROM owns_items oi WHERE oi.trainer_id IN
            (SELECT t.id
                FROM trainer t
                JOIN pokemon_creature p ON t.id = p.trainer
                WHERE p.id = NEW.creature)
        AND oi.item_id =
            (SELECT mm.id
                FROM move_machine mm
                JOIN move m ON mm.move = m.id
                WHERE m.id = NEW.move);

        --delete previous move
        DELETE FROM cause_or_receive WHERE slot = NEW.slot
        AND creature = NEW.creature AND move != NEW.move;

    ELSE
        --error message
        INSERT INTO warnings(affected_table, error_message, date, user_name)
        SELECT 'trainer', concat('<', t.class, '><', t.name,
        '> attempted to teach his <', s.id, '> the move <', m.name,
        '> without the necessary move machine.'), CURRENT_DATE, t.id
        FROM trainer t JOIN pokemon_creature pc ON pc.trainer = t.id
        JOIN species s ON pc.species = s.id
        JOIN move m ON m.id = NEW.move
        WHERE pc.id = NEW.creature;
    END IF;
END;
$$ LANGUAGE plpgsql;
```

```

--revert insert
DELETE FROM cause_or_receive WHERE id = NEW.id
AND move = NEW.move AND creature = NEW.creature;

END IF;
RETURN NEW;
END
$$ LANGUAGE plpgsql;

DROP TRIGGER IF EXISTS teach_new_move ON cause_or_receive CASCADE;
CREATE TRIGGER teach_new_move AFTER INSERT ON cause_or_receive
FOR EACH ROW
EXECUTE FUNCTION teach_move();

```

6.7.2 Explanation

We want to check if a trainer has the required move machine to teach a Pokémon a specific move. Therefore, we designed this trigger to be executed after a new *INSERT* was done to the **Cause_Or_Recieve** table, which links the **Pokemon_Creature** and **Move** entities. We were required to do it *AFTER* the insertion as in one of the options we had to revert the change (remove the inserted row).

Whenever the trigger gets activated, the function *teach_move()* is executed. This one has two main scenarios: the trainer has the necessary move machine to teach the move to their Pokémon, or if they haven't.

To check if we are on the first case, we use the *EXISTS* clause and a subquery to find out if the entered Pokémon's *id* corresponds to that of a trainer who has the required move machine. We do so by joining the **Trainer** and **Pokemon_Creature** entities and specifying in a *WHERE* clause that the Pokémons *id* has to be equal to the one just entered (using the *NEW* command). Then, we join the **Trainer** and **Move_Machine** entities (by connecting them through the **Owns_Item** table) and add as a second condition to the *WHERE* clause that forces the *move* the machine teaches to be the same as the one introduced. This query will return a row if the trainer does have the necessary move machine or nothing otherwise. If it returns a row, the function will execute the first *IF* statement. If not, it will execute the *ELSE*.

Once we have identified the trainer of the entered *Pokémon* to meet the requirements, we must do two operations. Firstly, we remove the move machine from the trainer's items by using the *DELETE FROM* command stating we want to eliminate information from the **Owns_Item** table. To specify which rows we want to eliminate, we add a *WHERE* clause in which we equal the *trainer_id* attribute of this table to the entered Pokémons *trainer* (which we find by means of a subquery by joining the **Trainer** and **Pokemon_Creature** entities and adding a *WHERE* clause to select only the Pokémon that has just been inserted) and also equal the *item_id* attribute to the move machine we want to delete (which we get by using a subquery that links the **Move** and **Move_Machine** entities and adds a *WHERE* condition to only get the move machine that teaches the newly entered move).

Also, when teaching a new move we want to remove the one that used to take up the same slot (if there was any). To do so we use the *DELETE FROM* command and specify (using a *WHERE* clause) we want to delete those rows from the **Cause_Or_Recieve** table

(linking **Pokemon_Creature** and **Move**) that have the same slot as the newly entered move, belong to the same Pokémon creature, but connect to a different move. This way, we will first insert a new move with the same *id* and then remove the one that has that same *id* but a different move from the one we just inserted.

On the other hand, if the trainer doesn't have the necessary move machine to teach the move, the function performs two actions. First, we create a new entry in the **Warnings** table by using the *INSERT INTO* command and selecting '*trainer*' as the affected table, concatenating a string with the desired message, adding the current date, and selecting the trainer's *id* as the user. To do so, we link the **Trainer** and **Pokemon_Creature** tables and specify in the *WHERE* clause that we want this Pokémon to be the newly inserted one (*NEW* command). This already allows us to get the trainer *class*, *name*, and *id* which we need for the message. However, to get the specie's *id* and the name of the move we join the tables **Species** (linking it to the Pokémon) and **Move** (stating we want this to be the newly inserted one). Also, we use the *concat()* method to get the desired *error_message* attribute to insert it.

Lastly, if the trainer didn't meet the requirements, we are asked to revert the change and delete the newly inserted row to the **Cause_Or_Receive** table. Hence, we use the *DELETE FROM* command combined with a *WHERE* clause to specify we want to delete the row that has the same *id* as the just introduced one.

6.7.3 Query validation

To validate this trigger, we will insert two new moves: one to a Pokémon whose trainer has the required move machine and another one that doesn't. Hence, we will see if our trigger works in both cases.

For example, if we check the following row of the **Move_Machine** table we can see how the move with *id* = 68 is taught by the machine with *id* = 322.

	id	move
	21	322

68

If we now check the **Owns_Item** table, the trainer with *id* = 0 owns that move machine.

trainer_id	item_id	obtention_method
0	322	REWARD

And by checking the **Pokemon_Creature** table we can see that the same trainer owns the Pokémon with *id* = 8.

id	species	nickname	trainer
1	8	387 Blue	0

Now if we check the moves that this Pokemon has in the **Cause_Or_Receive** we get the following:

creature	move	id
8	79	4
8	147	3
8	71	1
8	73	2

Now we try to insert the move with $id = 68$ on its first position by executing the following query:

```
INSERT INTO cause_or_receive(creature, move, id) VALUES (8, 68, 1);
```

If we now check the **Cause_Or_Receive** table again, we can see how this one was updated as this Pokémon's trainer had the necessary move machine to teach the move. Notice how the first move of the Pokémon ($id = 1$) now represents the move with $id = 68$.

creature	move	id
8	79	4
8	147	3
8	68	1
8	73	2

Also, if we check if the trainer with $id = 0$ still has the move machine by executing the following simple query, we can see how no rows are returned, meaning the move machine was correctly deleted from the trainer's items:

```
SELECT *
  FROM trainer t JOIN owns_items oi ON oi.trainer_id = t.id
    JOIN move_machine mm ON oi.item_id = mm.id
 WHERE t.id = 0 AND mm.id = 322;
```

t.id	class

As we can see, the operation was successful for a Pokémon whose trainer met the requirements. Now we will do the opposite case by performing the same query again (now that the trainer with $id = 0$ no longer has the move machine).

```
INSERT INTO cause_or_receive(creature, move, id) VALUES (8, 68, 2);
```

We can see how the **Cause_Or_Receive** table stayed the same (the move with $id = 68$ wasn't added on the second position or anywhere else, meaning it didn't replace the past move and its insertion was reverted).

creature	move	id
8	79	4
8	147	3
8	68	1
8	73	2

Also, if we check the warnings table we can see how we get the appropriate error message.

aff...	error_message	date	user_name
1	trainer <Expert><Gema> attempted to teach his <387> the move <counter> without the necessary move machine.	2023-05-20	0

6.8 Trigger 2

6.8.1 Solution

```
DROP FUNCTION IF EXISTS battle_reward CASCADE;
CREATE OR REPLACE FUNCTION battle_reward() RETURNS TRIGGER AS $$
BEGIN
    --Add winner gold and xp
    UPDATE trainer
        SET gold = gold + NEW.gold,
            experience = experience + 2.0/3.0 * NEW.xp
        FROM team tm
        WHERE tm.id = NEW.winner AND tm.trainer = trainer.id;

    --Add loser xp
    UPDATE trainer
        SET experience = experience + 1.0/3.0 * NEW.xp
        FROM team tm
        WHERE tm.id = NEW.loser AND tm.trainer = trainer.id;

    --Subtract gold if defeated by villain
    UPDATE trainer
        SET gold = gold - NEW.gold
        FROM evil_criminal ec
        JOIN team tec ON ec.id = tec.trainer
        JOIN team tt ON tt.id = NEW.loser
        WHERE tec.id = NEW.winner AND tt.trainer = trainer.id;

    --Give Gym Leader item to its defeater
    INSERT INTO owns_items(trainer_id, item_id, obtention_method, date_time)
        SELECT NEW.winner, i.id, 'WON IN BATTLE', CURRENT_TIMESTAMP
        FROM team t JOIN gym_leader gl ON t.trainer = gl.id
        JOIN gym g ON gl.gym = g.id JOIN item i ON g.id_item = i.id
        WHERE t.id = NEW.loser;

    RETURN NEW;
END
$$ LANGUAGE plpgsql;

DROP TRIGGER IF EXISTS battle_result ON battle CASCADE;
CREATE TRIGGER battle_result AFTER INSERT ON battle
FOR EACH ROW
EXECUTE FUNCTION battle_reward();
```

6.8.2 Explanation

This trigger was defined to make the result of battles consistent. Therefore, we designed it to be activated right after a new battle is inserted. When this happens, the function *battle_reward()* gets executed.

The function we designed has been split into four different tasks. Firstly, we *UPDATE* the **Trainer** entity by adding to the winner of the battle the *gold* reward and increase their *experience* by $\frac{2}{3}$ of the battle's *xp*. To do so, we just add to these attributes the newly inserted ones (using the *NEW* command). To identify which trainer we need to add it to, we do a *JOIN* with the **Team** entity and add a *WHERE* clause to specify that the row we want to update is that of the trainer that belongs to the team who was victorious.

Similarly, to add $\frac{1}{3}$ of the battle's XP to the loser we *UPDATE* the **Trainer** entity by adding the inserted attribute (*NEW*) to the one they had before. To identify which row we want to

change, we add a *WHERE* clause that links the trainer's team with the *loser* attribute of the newly inserted battle.

Moreover, if the *winner* of the battle was a criminal, the gold reward they receive must be subtracted from the other trainer. To do so, we *UPDATE* again the **Trainer** table and *SET* their *gold* attribute to its previous value minus the inserted one (*NEW*). To ensure we only do this whenever the winner is a criminal, we join the **Evil_Criminal** and **Team** entities and add a *WHERE* clause that the criminal's team has to be the same as the *winner* attribute of the inserted battle. This will only leave us with the battles in which a criminal was victorious. However, to specify which row of the defeated trainer we want to change, we join the entity **Team** once more (now linking it to the *loser* attribute of the new battle) and add a second condition to the *WHERE* clause that links the team to the **Trainer** table we are updating.

Lastly, if a Gym Leader is beaten, we want to give its item (the one on their gym) to its defeater. To do so, we use the *INSERT INTO* command to add a new row to the **Owns_Item** table (create a new relationship between a **Trainer** and an **Item**), in which we insert the *winner* attribute of the new battle (using the *NEW* command) as the owner of the item, specify the obtention method as '*WON IN BATTLE*', and add the current time as the *date_time* of obtention (using the *CURRENT_TIMESTAMP* method). Lastly, to get the *item* of the Gym Leader and add it to the table in the *item_id* column (to correctly link the trainer and the item) we need to do a set of *JOINS*. First, we join the **Team** and **Gym_Leader** entities to get the team of the Gym Leader, then we join the **Gym_Leader** to its **Gym** and lastly we link the **Gym** to the **Item** it has. To specify we want to get the Gym Leader who lost this battle we add a *WHERE* clause linking the leader's team with the *loser* of the newly inserted battle.

6.8.3 Query validation

To validate this trigger we will insert four different battles (one of each of the mentioned cases) and see their effects. First, we will look at four different trainers:

Regulat trainers and their teams (**Trainer** and **Team** entities): :

	id	class	name	experience	gold
1	0	Expert	Gema	536957	3820
2	1	Passionate Rider	Genie	915295	8092

	id	trainer
1	196	0
2	392	1

Evil criminal and their team (**Evil_Criminal**, **Trainer**, and **Team** entities):

	id	money_steal	signat...	is_leader	sidekick	criminal_organization
1	406	242	You want a to...	false	400	Rocket

	id	class	name	experience	gold
407	406	Rocket Executive	Lucien	330109	9634

	id	trainer
407	13	406

Gym Leader and their team (*Trainer*, *Gym_Leader*, and *Team* entities):

	id	class	name	experience	gold
505	504	Gym Leader	Koga	1365804	5981
	id	gym	id	trainer	
1	504	527	505	424	504

Now that we have correctly identified these four trainers we can proceed to insert the battles.

First, we will insert a regular battle in which one regular trainer beats another one.

```
INSERT INTO battle(id, winner, loser, gold, xp)
VALUES(600, 196, 392, 100, 100);
```

If we run the trigger, then do the insert above and check the tables we get the following:

Before:

	id	class	name	experience	gold
1	0	Expert	Gema	536957	3820
2	1	Passionate Rider	Genie	915295	8092

After:

	id	class	name	experience	gold
1	0	Expert	Gema	537024	3920
2	1	Passionate Rider	Genie	915328	8092

As you can see, the winner of the battle (trainer with *id* = 0) increased their gold by 100 units (*gold* attribute of inserted battle) while the other trainer stayed with the same amount. Also, notice how the experience of the first trainer got increased by 67 units ($\frac{2}{3}$ of 100, *experience* of the new battle) and the second trainer increased its experience by 33 units ($\frac{1}{3}$ of a 100).

If we now insert a new battle in which an evil criminal defeats a regular trainer we obtain the following:

```
INSERT INTO battle(id, winner, loser, gold, xp)
VALUES(601, 13, 392, 100, 100);
```

	id	class	name	experience	gold
2	1	Passionate Rider	Genie	915361	7992
407	406	Rocket Executive	Lucien	330176	9734

As you can see, the first trainer has decreased its gold by 100 units, while the second trainer has increased it by the same amount. Meaning that the evil criminal took the gold reward from the defeated trainer. Also if you check, the *experience* attribute has also been updated accordingly as the previous example.

Lastly, we insert a battle in which the defeated trainer is a Gym Leader and obtain the following:

```
INSERT INTO battle(id, winner, loser, gold, xp)
VALUES(603, 196, 424, 100, 100);
```

	id	class	name	experience	gold
1	0	Expert	Gema	537091	4020
505	504	Gym Leader	Koga	1365837	5981

The first thing we can see is that, as in the previous examples, the *gold* and *experience* attributes have been updated correctly (increasing *gold* by a 100 for the first trainer and staying the same for the second one; and increasing $\frac{2}{3}$ of 100 of experience for the winner and $\frac{1}{3}$ for the loser).

However, we should now check if the winner has obtained the corresponding item. If we check the previous images we can see how the gym of the Gym Leader with an *id* = 504 has an *id* of 527. From the table **Gym** we can see how this one gives the item with *id* = 353.

	id	medal	id_item	id_type
1	527	Sould Badge	353	4

Now if we check the **Owns_Item** table we can see there is a new entry linking the winner trainer (*id* = 196) with the gym's item (*id* = 353) with the designated obtention method.

	trainer_id	item_id	obtention_method	date_time
1	196	353	WON IN BATTLE	2023-05-20 08:01:16.565669

All of these inserts represent the different cases our trigger manages and prove its correct functioning by updating the trainers as expected.

7 Let's go Pokéshopping

7.1 Query 1

7.1.1 Solution

```
SELECT item_id, i.name, p.trainer_id, t.name
FROM purchase as p
JOIN item as i on p.item_id = i.id
JOIN trainer as t ON p.trainer_id = t.id
WHERE trainer_id = (SELECT p.trainer_id
FROM purchase as p
GROUP BY p.trainer_id
ORDER BY SUM(p.final_price * amount) DESC
LIMIT 1)
ORDER BY amount DESC
LIMIT 1;
```

	item_id	i.name	trainer_id	t.name
1	316	tm12	70	Freeman

7.1.2 Explanation

For this query, we have to filter by the trainer ID that has spent the most amount of money in stores. To do that, we have done a subquery that selects the trainer ID from the **Purchase** table group and does a group by **Trainer** so we could sum the **final_price** and amount by Trainer. In order to obtain the one that spent the most, we organized the result by the sum count in descendent order and limited by one. Finally, to select its favorite item we order by amount in descendent order and also limit by 1.

7.1.3 Query validation

	store_id	item_id	trainer_id	amount	final_price	discount	date
11	581	56	0	1	1000	0	2008-1
12	572	312	0	1	50000	0	2007-0
13	593	62	0	2	2000	0	2012-1
14	569	441	1	2	0	10	2011-1
15	564	555	1	2	0	56	1996-0
16	581	544	1	3	0	0	2019-0
17	561	252	1	2	4000	0	2003-1
18	592	256	1	3	4000	0	2003-0
19	590	66	1	2	20	0	2001-0
20	565	409	2	1	0	0	2017-1

If we select everything from the **Purchase** table and order by **trainer_id**, we obtain that for instance Trainer with ID = 1, has purchased a total of $8000 + 12000 + 40$. If we do the sum up just with the following query, we obtain the same result. This confirms obtaining the player that has spent the most amount of money in stores. The rest are simply joins that were validated in phase 2 and an order by amount in descending order.

```

SELECT p.trainer_id, SUM(p.final_price * amount) AS money
FROM purchase AS p
GROUP BY p.trainer_id
ORDER BY SUM(p.final_price * amount) DESC;

```

	trainer_id	money
1	0	78800
2	1	20040
3	2	8500
4	3	46420
5	4	152081
6	5	45363
7	6	56266

7.2 Query 2

7.2.1 Solution

```

SELECT o.trainer_id, t.name, t.experience, COUNT(DISTINCT
o.item_id)
FROM owns_items AS o
JOIN trainer AS t ON o.trainer_id = t.id
JOIN item AS i ON o.item_id = i.id
WHERE t.id IN (SELECT t.id
FROM trainer AS t
LEFT JOIN purchase AS p ON t.id = p.trainer_id
WHERE p.trainer_id IS NULL
ORDER BY t.name)
GROUP BY o.trainer_id, t.name, t.experience
HAVING COUNT(DISTINCT o.item_id) > 10 AND t.experience > 3000
ORDER BY COUNT(DISTINCT o.item_id);

```

trainer_id	name	experience	count
1	John	3500	12

7.2.2 Explanation

Since the query needs to list all the players that have never bought an item, we do a WHERE clause and a subquery to select them. To do so, we do a left join on the trainer ID. Finally, we do a group by ID, name and experience to filter those that have over 10 different items and experience is greater than 3000.

7.2.3 Query validation

```
SELECT trainer_id  
FROM purchase  
group by trainer_id;
```

A screenshot of a PostgreSQL query output window titled "Output" and "postgres.public". It shows a table with one column labeled "trainer_id" containing five rows: 501, 502, 503, 504, and 505. The values 326, 385, 371, 492, and 365 are displayed to the right of the column header.

trainer_id	
501	326
502	385
503	371
504	492
505	365

```
SELECT *  
FROM trainer;
```

A screenshot of a PostgreSQL query output window titled "Output" and "postgres.public.trainer". It shows a table with five columns: "id", "class", "name", "experience", and "gold". There are three rows with data: row 501 has class "Aqua Biologist", name "Magdalen", experience 1243163, and gold 4813; row 502 has class "Magma Grunt", name "Kenyetta", experience 212564, and gold 3490; row 503 has class "Magma Admin", name "Shara", experience 48403, and gold 1965.

	id	class	name	experience	gold
501	461	Aqua Biologist	Magdalen	1243163	4813
502	463	Magma Grunt	Kenyetta	212564	3490
503	462	Magma Admin	Shara	48403	1965

Since this query didn't return anything, we do a group by trainer_id in **Purchase** and compare the number of rows to the number of **Trainers**. In both cases, we obtain 537 rows meaning that all trainers stored in the system have at least purchased something.

7.3 Query 3

7.3.1 Solution

```
SELECT DISTINCT hf.flavour, COUNT(DISTINCT i.id) as num_flavours  
FROM sell as s  
JOIN item as i on s.item_id = i.id  
JOIN berrie as b on i.id = b.id_item  
JOIN has_flavour hf on b.id_item = hf.id_berrie  
JOIN flavour as f on hf.flavour = f.name  
GROUP BY hf.flavour  
ORDER BY COUNT(DISTINCT i.id)  
LIMIT 1;
```

A screenshot of a PostgreSQL query output window showing the results of the query. It has two columns: "flavour" and "num_flavours". The first row shows "flavour" as "bitter" and "num_flavours" as "19".

flavour	num_flavours
bitter	19

7.3.2 Explanation

To start, we do a join from **sell** with **item** since we want the items that can be purchased in stores and **berrie** because those are the only elements that can have flavors. Then, we join berrie with flavor through **has_flavour** to be able to access flavors information. We do a group by flavor, count the distinct IDs that appear and order by that in ascending order. Finally, we limit by 1 to have the least possible.

7.3.3 Query validation

To validate it, we have eliminated the LIMIT by 1, so all the flavors appeared together with the number of times used.

	flavour	num_flavours
1	bitter	19
2	spicy	21
3	sour	22
4	sweet	24
5	dry	27

For each flavor, we have done the following query in order to know the number of rows. In our case, dry gives us 27 rows, sour 22, spicy 21, sweet 24 and bitter 19. These results are equal to the ones obtained before.

Query for dry flavour:

```
SELECT DISTINCT i.id, i.name, hf.flavour, hf.strength
FROM sell as s
JOIN item as i on s.item_id = i.id
JOIN berrie as b on i.id = b.id_item
JOIN has_flavour hf on b.id_item = hf.id_berrie
JOIN flavour as f on hf.flavour = f.name
WHERE hf.flavour LIKE 'dry'
ORDER BY i.id;
```

Query for sour flavour:

```
SELECT DISTINCT i.id, i.name, hf.flavour, hf.strength
FROM sell as s
JOIN item as i on s.item_id = i.id
JOIN berrie as b on i.id = b.id_item
JOIN has_flavour hf on b.id_item = hf.id_berrie
JOIN flavour as f on hf.flavour = f.name
WHERE hf.flavour LIKE 'sour'
ORDER BY i.id;
```

Query for spice flavour:

```
SELECT DISTINCT i.id, i.name, hf.flavour, hf.strength
FROM sell as s
JOIN item as i on s.item_id = i.id
JOIN berrie as b on i.id = b.id_item
JOIN has_flavour hf on b.id_item = hf.id_berrie
JOIN flavour as f on hf.flavour = f.name
WHERE hf.flavour LIKE 'spicy'
ORDER BY i.id;
```

Query for sweet flavour:

```
SELECT DISTINCT i.id, i.name, hf.flavour, hf.strength
FROM sell as s
JOIN item as i on s.item_id = i.id
JOIN berrie as b on i.id = b.id_item
JOIN has_flavour hf on b.id_item = hf.id_berrie
JOIN flavour as f on hf.flavour = f.name
WHERE hf.flavour LIKE 'sweet'
ORDER BY i.id;
```

Last query for dry flavour:

```
SELECT DISTINCT i.id, i.name, hf.flavour, hf.strength
FROM sell as s
JOIN item as i on s.item_id = i.id
JOIN berrie as b on i.id = b.id_item
JOIN has_flavour hf on b.id_item = hf.id_berrie
JOIN flavour as f on hf.flavour = f.name
WHERE hf.flavour LIKE 'bitter'
ORDER BY i.id;
```

7.4 Query 4

7.4.1 Solution

```
SELECT combined.item_id, MAX(combined.trainers) AS trainers,
MAX(combined.stock) AS stock, MAX(combined.sum) AS sum
FROM
(
    SELECT item_id, trainers, 0 AS stock, 0 AS sum
    FROM
    (
        SELECT item_id, COUNT(DISTINCT trainer_id) AS trainers
        FROM owns_items
        GROUP BY item_id
    ) AS trainers_query
    UNION
```

```

SELECT item_id, 0 AS trainers, stock, 0 AS sum
FROM sell
WHERE stock = (SELECT MIN(stock) FROM sell)

UNION

SELECT item_id, 0 AS trainers, 0 AS stock, SUM(amount) AS sum
FROM purchase
GROUP BY item_id
) AS combined
GROUP BY combined.item_id
ORDER BY trainers, stock, sum
LIMIT 1;

```

	item_id	trainers	stock	sum
1	551	5	0	74

7.4.2 Explanation

This query aims to join three different queries with different criteria each. This is the reason why we have all combined with a UNION and from the result, we have selected the maximum value of all three. For example in stock when looking at the number of times each Item has been purchased, we fill this column with zeros, so at the end we want to have the maximum value which will correspond to the query that gives this solution.

7.4.3 Query validation

If we eliminate the LIMIT by 1, we will see that columns are combined and correspond to the three queries done separately.

	item_id	trainers	stock	sum
1	551	5	0	74
2	269	6	0	26
3	441	6	0	69
4	83	7	0	0
5	209	7	0	0
6	144	7	0	0
7	191	7	0	12
8	264	7	0	29
9	159	7	0	60

Looking at all the rows, we'll see that stock is always zero but that happens because we select the minimum stock from sell since the objective of this query is to select the rarest item and that value is zero.

```
--number of items owned by trainers
SELECT item_id, COUNT(DISTINCT trainer_id) as trainers
FROM owns_items
GROUP BY item_id
ORDER BY trainers;
```

The screenshot shows a database interface with a results grid. The columns are labeled 'item_id' and 'trainers'. The data consists of 8 rows:

	item_id	trainers
1	551	5
2	269	6
3	441	6
4	211	7
5	210	7
6	209	7
7	83	7
8	144	7

```
--store that each store has
SELECT *
FROM sell
WHERE stock = (SELECT MIN(stock)
FROM sell)
ORDER BY stock ;
```

Or

```
SELECT *
FROM sell
WHERE stock = (SELECT stock
FROM sell
GROUP BY stock
ORDER BY stock
LIMIT 1)
ORDER BY stock ;
```

The screenshot shows a database interface with a results grid. The columns are labeled 'store_id', 'item_id', 'stock', and 'discount'. The data consists of 9 rows:

	store_id	item_id	stock	discount
1	588	554	0	10
2	581	544	0	0
3	581	517	0	0
4	574	48	0	0
5	595	228	0	0
6	580	432	0	22
7	586	388	0	26
8	575	395	0	0
9	562	494	0	14

```
-- number of times each item has been purchased
SELECT item_id, SUM(amount) as sum
FROM purchase
GROUP BY item_id
ORDER BY sum;
```

The screenshot shows a table with two columns: 'item_id' and 'sum'. The data is as follows:

item_id	sum
447	551
448	552
449	553
450	554
451	555
452	556
453	557
454	558
455	559
456	560
457	561

Overall, we see that results match with the combined query.

7.5 Query 5

7.5.1 Solution

```
SELECT
    ( SUM(money) - SUM(total_discount) ) AS total_discount_paid
FROM (
    SELECT
        (i.base_price * p.amount) AS money,
        (p.final_price * p.amount) AS total_discount
    FROM purchase AS p
    JOIN item AS i ON i.id = p.item_id
    WHERE EXTRACT(YEAR FROM p.date_time) = 2020
) AS discounts;
```

The screenshot shows a table with one row and two columns. The first column is labeled 'total_discount_paid' and the value is 429937.

total_discount_paid
429937

7.5.2 Explanation

To give the total discounts, we have done a subquery with a WHERE clause to filter date_time by year 2020, and from that JOIN with **Item** to calculate the price that would be paid without the discount, multiplying the base_price times the amount for each **purchase**. Then, we do a similar operation to get the final_price with the discount. From those results, we sum all the total amount paid with and without the discount, subtract them and we get the total_discounts in that year.

7.5.3 Query validation

```
SELECT i.id, base_price, p.amount,
       (i.base_price * p.amount) AS money,
       p.final_price,
       (p.final_price * p.amount) AS discount
  FROM purchase AS p
 JOIN item AS i ON i.id = p.item_id
 WHERE EXTRACT(YEAR FROM p.date_time) = 2020;
```

	id	base_price	amount	money	final_price	discount
1	26	700	3	2100	315	945
2	110	2000	2	4000	2000	4000
3	492	0	10	0	0	0
4	123	50	2	100	50	100
5	213	1000	10	10000	1000	10000
6	107	3000	2	6000	3000	6000
7	282	1000	2	2000	1000	2000
8	79	400	2	800	348	696

This query helps us check that we properly obtain the money paid without and with the discount. For example, the item with ID 26 was three times purchased in the same date_time, so the amount paid will be equal to its base_price times the amount. Same for the discount but multiplying final_price instead of base_price.

7.6 Query 6

7.6.1 Solution 1

```
SELECT *
  FROM berrie
 WHERE berrie.id_item = (
  SELECT item_id
    FROM owns_items
   JOIN item i ON i.id = owns_items.item_id
   JOIN berrie b ON i.id = b.id_item
  WHERE obtention_method LIKE '%HARVESTED'
 GROUP BY item_id
 ORDER BY COUNT(obtention_method) DESC, item_id DESC
LIMIT 1);
```

	id_item	growing_time	maximum_amount	natural_gift_powder	size	softne
1	186	24	5	80	41	

	natural_gift_powder	size	softness	soil_dryness	firmness	extra_int
1	80	41	60	7	soft	61

7.6.2 Solution 2

```
SELECT      b.id_item,          growing_time,          maximum_amount,
natural_gift_powder,    size,    softness,    soil_dryness,    firmness,
extra_int
FROM owns_items oi
JOIN berrie b ON oi.item_id = b.id_item
WHERE oi.obtainion_method LIKE '%HARVESTED'
GROUP BY oi.item_id, b.id_item, growing_time, maximum_amount,
natural_gift_powder, size, softness, soil_dryness, firmness,
extra_int
ORDER BY COUNT(DISTINCT oi.trainer_id) DESC, item_id DESC
LIMIT 1;
```

	id_item	growing_time	maximum_amount	natural_gift_powder	size	softness
1	186	24	5	80	41	41

	natural_gift_powder	size	softness	soil_dryness	firmness	extra_int
1	80	41	60	7	soft	61

7.6.3 Explanation

This last query has two approaches. On the first one, there's a subquery that searches the id of the **berrie** that has been harvested most times. We obtain this with a WHERE clause to filter the obtainion method of **owns_items** to those that contain "HARVESTED". Then, it is grouped by item ID for the counting and also ordered by item to be able to obtain the same results in both approaches since it is limited by 1.

Second approach eliminates the subquery and adds attributes to consider when printing a **berrie** in the group by. Also, instead of counting harvested, we count DISTINCT trainer_id because we still have a WHERE clause that selects those obtainion_method that contain HARVESTED.

7.6.4 Query validation

```
SELECT oi.item_id, COUNT(DISTINCT oi.trainer_id)
FROM owns_items oi
JOIN berrie b ON oi.item_id = b.id_item
WHERE oi.obtainion_method LIKE '%HARVESTED'
GROUP BY oi.item_id
ORDER BY item_id;
```

To validate it, we have done this query that does the group by item_id and filters by those containing harvested in obtainion_method. It counts DISTINCT trainer_ids and it is ordered by item_id to ease the work.

To validate the counting, we have compiled the following query.

```
SELECT *
FROM owns_items
JOIN berrie b ON item_id = b.id_item
WHERE obtention_method LIKE '%HARVESTED'
ORDER BY item_id ;
```

As we can observe item with ID 126 has been harvested three times, so this confirms the three obtained in the first query, which at the end is signaling that if it's ordered by that count and limited by 1, we will obtain the berry that has been harvested the most times.

7.7 Trigger 1

7.7.1 Solution

```
DROP FUNCTION IF EXISTS calculate_healing CASCADE;
CREATE OR REPLACE FUNCTION calculate_healing() RETURNS TRIGGER AS
$$
DECLARE
    heal INTEGER;
    maxHeal INTEGER;
    currentHP INTEGER;
BEGIN
    IF OLD.item_id IN (SELECT id From healing_item) THEN
        currentHP := (SELECT hp FROM pokemon_creature
```

```

        WHERE pokemon_creature.id = (SELECT p.id FROM
pokemon_creature AS p
        WHERE p.position = 1 AND p.trainer = OLD.trainer_id));

heal := currentHP + (SELECT hi.status_heal
FROM healing_item AS hi
WHERE OLD.item_id = hi.id);

maxHeal := (SELECT ((2 * hp + (level ^ 2) / 4) / 100 + 5) *
nature
FROM pokemon_creature
WHERE pokemon_creature.id = (SELECT p.id FROM
pokemon_creature AS p
WHERE p.position = 1 AND p.trainer = OLD.trainer_id));

IF heal <= maxHeal THEN
    UPDATE pokemon_creature
    SET hp = heal
    WHERE pokemon_creature.id = (SELECT p.id
FROM pokemon_creature AS p
WHERE p.position = 1 AND p.trainer = OLD.trainer_id);
ELSE
    UPDATE pokemon_creature
    SET hp = maxHeal
    WHERE pokemon_creature.id = (SELECT p.id
FROM pokemon_creature AS p
WHERE p.position = 1 AND p.trainer = OLD.trainer_id);
END IF;

RETURN NEW;
END
$$ LANGUAGE plpgsql;

DROP TRIGGER IF EXISTS healing_result ON owns_items CASCADE;
CREATE TRIGGER healing_result AFTER DELETE ON owns_items
FOR EACH ROW
EXECUTE FUNCTION calculate_healing();

```

7.7.2 Explanation

Whenever a delete is done in `owns_items`, we check whether this `OLD.id` was a healing item by doing a subquery. If it is, we set the current HP by first finding the `pokemon_creature` that is in the first position and whose trainer was the one we have just deleted. Then, we want to sum up that value plus the status heal of the Item we have removed specifying this in the `WHERE` clause. In order not to surpass the maximum hp we calculate it with the formula for the same `pokemon_creature` as before. If the sum is less than the maximum, we update this hp of the `pokemon_creature` to the sum, otherwise we set it to the maximum Heal calculated.

7.7.3 Query validation

For the query validation, whenever I delete an item I want to make sure I am getting it status_heal and maxHeal correctly, so I insert those values into the warnings table.

```
INSERT INTO warnings
SELECT hi.status_heal
FROM healing_item AS hi
WHERE OLD.item_id = hi.id;
```

Here are the tests done to test the trigger.

```
DELETE FROM owns_items
WHERE item_id = 26 AND trainer_id = 4;
```

For item_id 26, we found in healing_item that status_heal is 77 and calculating with the level, base stat and nature of the pokemon_creature we obtain maxHeal 117 which corresponds to the **warnings** table. Here it's important to note that the pokemon_creature chosen is Noel since it has trainer_id = 4 and position =1.

The screenshot shows three tables in MySQL Workbench:

- healing_item**: Contains rows for item_id 32 (hp_recovered 80, status_heal 28) and item_id 26 (hp_recovered 50, status_heal 77).
- warnings**: Contains two rows with message values 77 and 117.
- pokemon_creature**: Contains a single row for Noel (id 426, hp 120, maxHeal 117, trainer_id 4).

Finally, we'll check whether it has updated or not the hp. In our trigger, we modify Noel's row in pokemon_creature and set his hp to 83, which was previously defined as 6.

The screenshot shows the **pokemon_creature** table with one row for Noel (id 426, hp 83, maxHeal 117, trainer_id 4).

7.8 Trigger 2

7.8.1 Solution

```
DROP TABLE IF EXISTS warnings CASCADE;
CREATE TABLE warnings (
```

```

    message TEXT
) ;

DROP FUNCTION IF EXISTS modify_purchase CASCADE;
CREATE OR REPLACE FUNCTION modify_purchase() RETURNS TRIGGER AS $$

DECLARE
    price INTEGER;
    newStock INTEGER;
    newGold INTEGER;
BEGIN

    price := (SELECT i.base_price
              FROM item as i
              WHERE NEW.item_id = i.id) -
        (SELECT i.base_price
              FROM item as i
              WHERE NEW.item_id = i.id) * NEW.discount / 100;

    UPDATE purchase
    SET final_price = price
    WHERE NEW.trainer_id = trainer_id AND NEW.item_id = item_id AND
        NEW.store_id = store_id AND NEW.date_time = date_time;

    newGold := (SELECT gold FROM trainer WHERE NEW.trainer_id =
trainer.id) - (price * NEW.amount);
    IF newGold < 0 THEN
        INSERT INTO warnings
        SELECT CONCAT(
            'Trainer #', t.id,
            ' tried to buy ', i.name,
            ' but his card has been declined, insufficient funds.')
    END IF;

    FROM purchase
    JOIN trainer as t ON purchase.trainer_id = t.id
    JOIN item as i ON purchase.item_id = i.id
    JOIN shop as s ON purchase.store_id = s.id
        WHERE NEW.trainer_id = t.id and NEW.item_id = i.id AND
NEW.store_id=s.id and NEW.date_time = date_time;

    DELETE FROM purchase
    WHERE trainer_id = NEW.trainer_id and item_id = NEW.item_id
and store_id = NEW.store_id and
date_time = NEW.date_time and amount=NEW.amount and
discount=NEW.discount;
    ELSE

        newStock = (SELECT stock FROM SELL
                    WHERE NEW.store_id = sell.store_id and NEW.item_id =
sell.item_id) - NEW.AMOUNT;
        IF newStock < 0 THEN

```

```

        INSERT INTO warnings
        SELECT CONCAT(
            'Trainer #', t.id,
            ' tried to buy ', i.name,
            ' at ', s.name, ' store in ',
            sub.name, ', however the store ran out of stock.'
        )
        FROM purchase
        JOIN trainer as t ON purchase.trainer_id = t.id
        JOIN item as i ON purchase.item_id = i.id
        JOIN shop as s ON purchase.store_id = s.id
        JOIN subarea as sub ON s.id = sub.id
        WHERE NEW.trainer_id = t.id and NEW.item_id = i.id and
NEW.store_id = s.id;

        DELETE FROM purchase
            WHERE trainer_id = NEW.trainer_id and item_id =
NEW.item_id and store_id = NEW.store_id and
            date_time = NEW.date_time and amount=NEW.amount and
discount=NEW.discount;

    else

        INSERT INTO owns_items
            VALUES (NEW.trainer_id, NEW.item_id, 'PURCHASED',
NEW.date_time);
        UPDATE trainer
        SET gold = newGold
        WHERE NEW.trainer_id = trainer.id;
        UPDATE sell
        SET stock = sell.stock - NEW.amount
        WHERE NEW.store_id = sell.store_id and NEW.item_id =
sell.item_id ;
        end if;

    end if;

    RETURN NEW;
END
$$ LANGUAGE plpgsql;

DROP TRIGGER IF EXISTS insert_purchase ON purchase CASCADE;
CREATE TRIGGER insert_purchase AFTER INSERT ON purchase
FOR EACH ROW
EXECUTE FUNCTION modify_purchase();

```

7.8.2 Explanation

After an insert on the **purchase** table, we calculate the final_price by applying the discount thanks to the base_price of the Item and the amount the trainer has purchased it. Then, we update this value in the corresponding row with the new trainer_id, store_id, item_id and date_time. As a third step, we calculate the remaining gold the trainer will have.

If it's negative, it means the trainer cannot buy this, so we insert a message indicating this into the **warnings** table and delete that row from the **purchase** table. Otherwise, we calculate the remaining stock by computing the difference between the stock for a specific item and store, and the amount of items purchased. Again, if it's negative, we insert the corresponding message into **warnings** and delete that row in **purchase**. Otherwise, it will mean that we need to create an entry in **owns_items**, update trainer's gold and also store's stock.

7.8.3 Query validation

First case: Trainer doesn't have enough gold to do this purchase.

```
INSERT INTO purchase(store_id, item_id, trainer_id, amount,
discount, date_time)
VALUES (582, 367, 407, 1, 30, '2000-01-01');
```

Before doing the insert, Trainer with ID 407 has 2084 gold.

130	407	Magma Geologist	Brenton	823483	2084
-----	-----	-----------------	---------	--------	------

A part of the purchase table looks like this.

4739	582	347	434	3	1000
4740	582	347	446	2	660
4741	582	347	459	2	1000
4742	582	367	104	2	50000
4743	582	367	106	1	39500
4744	582	367	188	2	50000
4745	582	367	224	2	40000
4746	582	367	273	6	50000
4747	582	367	407	9	50000
4748	582	367	457	2	50000

In this case, he wants to buy the item with ID 367 that costs 50000, so in this case the trainer has not enough gold to buy it.

201	367	tm63	50000	Teaches Embargo to a compat...	fa
-----	-----	------	-------	--------------------------------	----

As a result, we see that it is not added in the purchase table.

4739	582	347	434	3	1000
4740	582	347	446	2	660
4741	582	347	459	2	1000
4742	582	367	104	2	50000
4743	582	367	106	1	39500
4744	582	367	188	2	50000
4745	582	367	224	2	40000
4746	582	367	273	6	50000
4747	582	367	407	9	50000
4748	582	367	457	2	50000

Also, it has appeared the message in the **warnings** table.

message
1 Trainer #407 tried to buy tm63 but his card has been declined, insufficient fu...

Second case: Trainer has enough gold but there's not enough stock

```
INSERT INTO purchase(store_id, item_id, trainer_id, amount, discount, date_time)
VALUES (561, 67, 230, 2, 0, '2000-01-03');
```

Trainer's gold is 8921.

230	Cowgirl	Ashlee	936598	8921
-----	---------	--------	--------	------

Item's stock in store 561 is 1.

561	67	1	0
-----	----	---	---

In the **warnings** table, it appears the corresponding message and if we want to look for this row in **purchase**, it doesn't exist.

message
1 Trainer #407 tried to buy tm63 but his card has been declined, insufficient fu...
2 Trainer #230 tried to buy red flute at Koelpin, Lindgren and Hintz store in Koelpin, Lindgren and Hintz, however

438	561	67	137	1	20
439	561	67	367	5	20

Third case: Trainer has gold and there's enough stock

```
INSERT INTO purchase(store_id, item_id, trainer_id, amount, discount, date_time)
VALUES (562, 320, 229, 4, 10, '2000-01-01');
```

Trainer gold is 11460:

8	229 Guy	Vernie	840327	11460
---	---------	--------	--------	-------

Store stock is 100:

1	562	320	100	0
---	-----	-----	-----	---

Base price of the item is 1000.

248	320 tm16	1000 Teaches Light Screen to a c...	f
-----	----------	-------------------------------------	---

	store_id	item_id	trainer_id	amount	final_price
686	562	320		120	1000
687	562	320		127	820
688	562	320		219	1000
689	562	320	229	4	900
689	229	4		10	2000-01-01 00:00:00.000000

As we can see, the purchase is added in the corresponding table, and the final_price corresponds to the 10% discount of the base_price of the Item with ID 320.

Then, in **trainers** gold is updated to 7860 , and in **sell** the stock to 96.

308	229 Guy	Vernie	840327	7860
-----	---------	--------	--------	------

1	562	320	96	0
---	-----	-----	----	---

8 Time to Explore

8.1 Query 1

8.1.1 Solution 1

```
SELECT region.name AS region_name, COUNT(subarea.id) AS subarea_count
FROM region
JOIN area ON region.name = area.name_region
JOIN subarea ON area.id = subarea.id_area
GROUP BY region.name
ORDER BY subarea_count DESC
LIMIT 1;
```

	region_name	subarea_count
1	sinnoh	174

8.1.2 Solution 2

```
SELECT region.name AS region_name, subarea_count
FROM region
JOIN (
    SELECT area.name_region, COUNT(subarea.id) AS subarea_count
    FROM area
    JOIN subarea ON area.id = subarea.id_area
    GROUP BY area.name_region
) AS subarea_counts ON region.name = subarea_counts.name_region
ORDER BY subarea_count DESC
LIMIT 1;
```

	region_name	subarea_count
1	sinnoh	174

8.1.3 Explanation

In the case of the first query, it selects the name from the **Region** table and counts the number of **Subareas** using the COUNT function. It then joins the **region** table with the **Area** table based on the *name_region* column and it joins the **area** table with the **subarea** table based on the *id* and *id_area* columns, respectively. The result is grouped by **region.name** in order to make sure that the COUNT() function calculates the amount of **subareas** for each unique region. The ORDER BY clause orders the result in descending order based on the count of **subareas** so that, finally, the LIMIT 1 ensures that only the **region** with the highest count is returned.

As for the second query, it achieves the same result by using a subquery to calculate the count of **subareas** for each **area** before joining with the **region** table. The subquery in the JOIN clause calculates the count of **subareas** for each area and associates it with the *name_region* column.

8.1.4 Query validation

The queries use some simple JOINS based on their relationships as location. The GROUP BY clause makes sure that all of the locations are specifically related to each **region**. In addition, by performing a similar query we can see all of the different **subareas** related to each **region** and we can make sure that Sinnoh is the one with the most instances of **subareas**.

	region_name	subarea_count
1	sinnoh	174
2	kanto	166
3	hoenn	142
4	johto	118

8.2 Query 2

8.2.1 Solution

```
SELECT s.name AS gym_name, a.name AS city_it_belongs_to,
       ar.name AS route_directly_connected, oa.name AS area_on_the_other_side
FROM subarea s
JOIN gym g ON s.id = g.id
JOIN area a ON s.id_area = a.id
JOIN city c ON a.id = c.id
JOIN connect_co ON a.id = co.id_area
JOIN route r ON co.id_route = r.id
JOIN area ar ON ar.id = r.id
JOIN connect_co2 ON r.id = co2.id_route
JOIN area oa ON oa.id = co2.id_area
WHERE a.name <> oa.name;
```

gym_name	city_it_belongs_to	route_directly_connected	area_on_the_other_side
6th Indigo League Gym	saffron city	route 5	cerulean city
8th Sinnoh League Gym	sunyshore city	route 223	sinnoh pokemon league
4th Sinnoh League Gym	pastoria city	route 212	hearthome city
1st Johto League Gym	violet city	route 31	dark cave
3rd Indigo League Gym	vermillion city	route 6	saffron city
2nd Indigo League Gym	cerulean city	route 24	route 25
2nd Sinnoh League Gym	eterna city	route 205	eterna forest
1st Hoenn League Gym	rustboro city	route 115	meteor falls
3rd Johto League Gym	goldenrod city	route 35	national park
4th Sinnoh League Gym	pastoria city	route 213	valor lakefront
2nd Sinnoh League Gym	eterna city	route 211	mt coronet
6th Johto League Gym	olivine city	route 39	route 38
3rd Hoenn League Gym	mauville city	route 111	route 113
1st Johto League Gym	violet city	route 36	route 37
5th Sinnoh League Gym	hearthome city	route 209	solaceon town
3rd Hoenn League Gym	mauville city	route 118	route 119
7th Indigo League Gym	cinnabar island	route 21	pallet town
1st Sinnoh League Gym	oreburgh city	route 207	route 206
2nd Hoenn League Gym	dewford town	route 106	route 105
4th Johto League Gym	ecruiteak city	route 42	mt mortar

8.2.2 Explanation

What this query has in particular is that we need to do several JOINS including three times with the **Area** table. Since the **Gyms** are a generalization of **Subareas**, we need to connect the table **Subarea** with the **Area** table and this one with the **City** table in order to get the city of each **gym**. In order to get the **Route** of each **City** we need to do several JOINS including the table **Connect_** and another one with area (since **Routes** are a generalization of **Areas**). Finally, so as to get the **areas** to the other side of the **routes** we need to do another connection with the tables **Connect_** and **Area**. As for the WHERE

clause, it ensures that the area name is not the same as the other side area name and avoids duplications.

8.2.3 Query validation

In order to validate this query we can take a look at the information given to us and corroborate some of the information. For example:

Route 104	Rustboro City	Petalburg City		Route 105	Grass
Route 105	Route 104			Route 106	Water
Route 106	Route 105			Dewford Town	Water
Route 107		Route 108	Dewford Town		Water
Route 108		Route 109	Route 107		Water
Route 109	Slateport City		Route 108		Water
Route 110	Mauville City	Route 103		Slateport City	Grass
Route 111	Route 113		Route 112	Mauville City	Sand
Route 112		Route 111	Lavaridge Town		Sand
Route 113		Route 111	Fallabor Town		Mud
Route 114		Fallabor Town		Meteor Falls	Rocks
Route 115	Meteor Falls			Rustboro City	Gravel
Route 116		Rusturf Tunnel	Rustboro City		Grass
Route 117		Mauville City	Verdanturf Town		Snow

The 1st Hoenn League Gym is placed in Rustboro City. That city would be directly connected through a Route to Petalburg City, Route 105, Meteor Falls and Rusturf Tunnel. This is exactly what we get as a result from the query:

#	gym_name	city_it_belongs_to	route_directly_connected	area_on_the_other_side
1	1st Hoenn League Gym	rustboro city	route 115	meteor falls
2	1st Hoenn League Gym	rustboro city	route 104	petalburg city
3	1st Hoenn League Gym	rustboro city	route 116	rusturf tunnel
4	1st Hoenn League Gym	rustboro city	route 104	route 105

8.3 Query 3

8.3.1 Solution

```

SELECT e.id, p.name AS pokemon_name, s.name AS subarea_name, e.meeting_method,
f.condition_type, f.condition_value, f.probability,
    f.min_level, f.max_level
FROM encounter e
JOIN subarea s ON e.subarea = s.id
JOIN area a ON s.id_area = a.id
JOIN region r ON a.name_region = r.name
JOIN find f ON f.id_encounter = e.id
JOIN species p ON f.id_species = p.id
WHERE r.name = 'sinnoh' AND e.meeting_method = 'surf';

```

	id	pokemon_name	subarea_name	meeting_method	condition_type	condition_value	probabilit
1	1021	zubat	b1f	surf			
2	1021	zubat	b1f	surf			
3	1010	zubat	4f	surf			
4	1010	zubat	4f	surf			
5	1001	zubat	mt coronet 1f route 207	surf			
6	1001	zubat	mt coronet 1f route 207	surf			
7	1093	zubat	b1f	surf			
8	1087	zubat	ravaged path area	surf			
9	1021	golbat	b1f	surf			
10	1021	golbat	b1f	surf			
11	1021	golbat	b1f	surf			
12	1010	golbat	4f	surf			
13	1010	golbat	4f	surf			
14	1010	golbat	4f	surf			
15	1001	golbat	mt coronet 1f route 207	surf			
16	1001	golbat	mt coronet 1f route 207	surf			
17	1001	golbat	mt coronet 1f route 207	surf			
18	1093	golbat	b1f	surf			
19	1087	golbat	ravaged path area	surf			
20	1080	golbat	inside h1f	surf			

The query result is too long to be able to show it completely in this report format.

8.3.2 Explanation

The query uses several JOINS in order to relate the information of the **encounters** (which is separated between the table **Encounter** and the table **Find**) with the **region**. It joins the **encounter** table with the **subarea** table based on the id of the **subarea**. It further joins the **subarea** table with the **area** table based on the id of the **area**. It finally reaches the **region** by joining this last table with **Region** using the region name. Then it joins the **find** table with the **encounter** table based on the id of the encounter and the **species** of the pokemons based on their id as well. The WHERE clause filters the result to only include encounters in the 'Sinnoh' region (**r.name = 'sinnoh'**) and with the meeting method set as 'surf' (**e.meeting_method = 'surf'**).

8.3.3 Query validation

We can use the canalave city area (subarea) as an example. Its id is the number 371. It has several surf encounters such as the following:

pelipper	371	surf			4	20	40
tentacool	371	surf			60	20	30
tentacruel	371	surf			5	20	40

We can also safely say that it belongs to the Sinnoh region.

sinnoh	canalave city	canalave city area	371
--------	---------------	--------------------	-----

Finally, we can see that all of this information is included in the result of the query:

	id	pokemon_name	subarea_name	meeting_method
1	967	tentacool	canalave city area	surf
2	967	tentacruel	canalave city area	surf
3	967	wingull	canalave city area	surf
4	967	pelipper	canalave city area	surf
5	967	pelipper	canalave city area	surf
6	967	shellos	canalave city area	surf
7	967	gastrodon	canalave city area	surf

8.4 Query 4

8.4.1 Solution

```
SELECT e.id, p.name AS pokemon_name, f.probability, f.min_level
FROM encounter e
JOIN find f ON e.id = f.id_encounter
JOIN species p ON f.id_species = p.id
WHERE f.probability <= 1 AND f.min_level =
(SELECT f2.min_level FROM find f2 WHERE f2.probability <= 1
ORDER BY min_level ASC LIMIT 1);
```

	id	pokemon_name	probability	min_level
1	1160	nidoran-m	1	2
2	1160	nidoran-m	1	2
3	1161	sentret	1	2
4	1161	sentret	1	2
5	779	wingull	1	2
6	1161	starly	1	2
7	1160	starly	1	2
8	1161	bidoof	1	2
9	1160	bidoof	1	2

8.4.2 Explanation

We first need to join the tables containing all the information of the **encounters** with the **species** to be able to show the pokémon's *names*. The WHERE clause filters those results based on the ones that have a *probability* of 1 or lower and also that have a *min_level* equal to the result of the subquery. The subquery finds the *minimum level* out of the encounters with 1% chance or lower.

8.4.3 Query validation

We can directly see the information in the data provided and we can see that all of the possible results are included in the query:

pokemon	subarealID	method	condition_	condition_value	chance%	min_level	max_level
togepi	373	gift			100	1	1
togepi	232	gift			100	1	1
mareep	233	gift			100	1	1
wooper	233	gift			100	1	1
slugma	233	gift			100	1	1
happiny	518	gift			100	1	1
riolu	455	gift			100	1	1
nidoran-m	474	walk	radar	on	1	2	2
nidoran-m	474	walk	radar	on	1	2	2
sentret	475	walk	radar	on	1	2	2
sentret	475	walk	radar	on	1	2	2
wingull	300	walk			1	2	2
starly	474	walk	radar	off	1	2	2
starly	475	walk	radar	off	1	2	2
bidoof	474	walk	radar	off	1	2	2
bidoof	475	walk	radar	off	1	2	2
growlithe	474	walk	item	pp up	4	2	2
growlithe	474	walk	item	sun stone	4	2	2
growlithe	475	walk	item	potion	4	2	2
growlithe	475	walk	item	twisted spoon	4	2	2

8.5 Query 5

8.5.1 Solution

```

UPDATE encounter
SET meeting_method =
CASE
    WHEN meeting_method = 'good' THEN 'good rod'
    WHEN meeting_method = 'old' THEN 'old rod'
    WHEN meeting_method = 'super' THEN 'super rod'
END
WHERE meeting_method IN ('good', 'old', 'super');

CREATE TEMPORARY TABLE available_methods AS
SELECT DISTINCT meeting_method
FROM encounter
WHERE meeting_method <> 'rock';

UPDATE encounter
SET meeting_method = (
    SELECT meeting_method
    FROM available_methods
    ORDER BY RANDOM()
    LIMIT 1
)
WHERE meeting_method = 'rock';

```

	id	subarea	meeting_method
122	268	104	old rod
123	269	104	super rod
124	284	117	good rod
125	285	117	old rod
126	287	117	super rod
127	290	118	good rod
128	291	118	old rod
129	292	118	super rod
130	295	119	good rod
131	296	119	old rod
132	297	119	super rod
133	300	120	good rod
134	301	120	old rod
135	302	120	super rod
136	306	122	good rod
137	307	122	old rod
138	308	122	super rod
139	311	123	good rod
140	312	123	old rod
141	313	123	super rod

And if we search based on the meeting_method 'rock' we would get 0 results.

8.5.2 Explanation

First what we simply do is replace all of the mentions of 'old', 'good' or 'super' with 'old rod', 'good rod' and 'super rod' respectively by creating an UPDATE that takes into account all of those cases. Then, we create a temporary table which contains all of the *meeting_methods* except 'rock'. This table is then used in the subquery of the second UPDATE in order to randomly select one of the other methods and insert it wherever the method 'rock' appears.

8.5.3 Query validation

Firstly, we can see that if we do the following query:

```
SELECT * FROM encounter WHERE meeting_method = 'rock';
```

We get 0 results.

<	<	0 rows	v	>	>I		↻	Q	■	+	-	↶	↷	↑	Tx: Auto	v

Secondly, we can also take a look at some of the instances in which the 'rock' method was used and see that they have all now changed to the same new method.

	id	subarea	meeting_method
1	34	14	rock
2	156	51	rock
3	162	52	rock
4	169	53	rock

	id	subarea	meeting_method
1	34	14	walk
2	156	51	walk
3	162	52	walk
4	169	53	walk

8.6 Query 6

8.6.1 Solution

```
SELECT DISTINCT a.name, r.pavement, f.condition_value
FROM area a
JOIN subarea s ON a.id = s.id_area
JOIN encounter e ON e.subarea = s.id
JOIN find f ON e.id = f.id_encounter
JOIN route r ON a.id = r.id
WHERE r.pavement = 'Grass' AND (f.condition_value like 'rain' OR
f.condition_value like 'night');
```

	name	pavement	condition_value
1	route 1	Grass	night
2	route 14	Grass	night
3	route 16	Grass	night
4	route 201	Grass	night
5	route 204	Grass	night
6	route 209	Grass	night
7	route 212	Grass	night
8	route 26	Grass	night
9	route 28	Grass	night
10	route 34	Grass	night
11	route 39	Grass	night
12	route 47	Grass	night
13	route 47	Grass	rain
14	route 5	Grass	night
15	route 7	Grass	night

8.6.2 Explanation

In this query we use the DISTINCT clause since there we would have had repeated values otherwise. We then use several JOINS to be able to relate all the information from the **encounters** to the **routes**. The WHERE clauses make sure that we only take into consideration those **Routes** that have 'grass' **pavement** and only those encounters that happen at night or while it is raining.

8.6.3 Query validation

We can see that there are a couple of instances with the id_encounter 553 that happen during rainy conditions:

47	9605	553	223	weather	rain
48	6926	553	129	weather	rain

This encounters happen in a route where the pavement is snow:

	id	pavement
1	348	Snow

However, if we were to change it to Grass we would be able to see that those encounters are added to our results (although only one result is added since we don't want repeated results)

```
UPDATE route
SET pavement = 'Grass'
WHERE id = 348;
```

	name	pavement	condition_value
1	route 1	Grass	night
2	route 14	Grass	night
3	route 16	Grass	night
4	route 201	Grass	night
5	route 204	Grass	night
6	route 209	Grass	night
7	route 212	Grass	night
8	route 26	Grass	night
9	route 28	Grass	night
10	route 34	Grass	night
11	route 39	Grass	night
12	route 44	Grass	rain
13	route 47	Grass	night
14	route 47	Grass	rain
15	route 5	Grass	night
16	route 7	Grass	night

8.7 Trigger 1

8.7.1 Solution

```
CREATE OR REPLACE FUNCTION add_city_info_function()
RETURNS TRIGGER AS $$
DECLARE
    region_name VARCHAR(255);
    gym_name VARCHAR(255);
    gym_leader_name VARCHAR(255);
    gym_leader_id INT;
    experience INT;
    gold INT;
    gym_medal VARCHAR(255);
    item_id INT;
    type_id INT;
    city_name VARCHAR(255);
```

```

    subarea_id INT;
BEGIN
-- Check if the region Lasalia exists
SELECT name INTO region_name FROM region WHERE name = 'Lasalia';

-- If not, create it
IF region_name IS NULL THEN
    INSERT INTO region (name) VALUES ('Lasalia') RETURNING name INTO region_name;
END IF;

-- Generate random values
city_name := 'City ' || (floor(random() * 1000) + 1)::text;
gym_name := (floor(random() * 1000) + 1)::text || ' Gym';
gym_medal := (floor(random() * 1000) + 1)::text || ' Badge';
gym_leader_name := 'Gym Leader ' || (floor(random() * 1000) + 1)::text;
experience := floor(random() * 1000) + 1;
gold := floor(random() * 10000) + 1;

-- Random item
SELECT id FROM item ORDER BY random() LIMIT 1 INTO item_id;

-- Random type
SELECT id FROM type ORDER BY random() LIMIT 1 INTO type_id;

-- Insert the city into area
INSERT INTO area (id, name, name_region) VALUES (NEW.id, city_name, region_name);

-- Get the next value from the subareas
DROP SEQUENCE IF EXISTS sub_area_id_seq;
CREATE SEQUENCE sub_area_id_seq;
PERFORM setval('sub_area_id_seq', (SELECT MAX(id) FROM subarea));
SELECT nextval('sub_area_id_seq') INTO subarea_id;

-- Create the subarea associated with the gym
INSERT INTO subarea (id, name, id_area) VALUES (subarea_id, gym_name, NEW.id);

-- Create the gym associated with the city
INSERT INTO gym (id, medal, id_item, id_type)
VALUES (subarea_id, gym_medal, item_id, type_id);

-- Get the next value from the trainers
DROP SEQUENCE IF EXISTS trainer_id_seq;
CREATE SEQUENCE trainer_id_seq;
PERFORM setval('trainer_id_seq', (SELECT MAX(id) FROM trainer));
SELECT nextval('trainer_id_seq') INTO gym_leader_id;

-- Insert the gym leader into the trainer table
INSERT INTO trainer (id, class, name, experience, gold) VALUES
(gym_leader_id, 'Gym Leader', gym_leader_name, experience, gold);

-- Create the gym leader associated with the gym
INSERT INTO gym_leader (id, gym) VALUES (gym_leader_id, subarea_id);

RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER add_city_info_trigger
BEFORE INSERT ON city
FOR EACH ROW
EXECUTE FUNCTION add_city_info_function();

```

8.7.2 Explanation

The trigger first checks whether the Lasalia **region** already exists, if it does not, it creates it. Then, we generate random values for the city name and everything related to the **Gym** and to the **Gym_leader**. These will compose the full information regarding the **city** in the different tables. Later, we have to start inserting into each corresponding table but it has to be in order so as not to mess up any foreign keys. We first, insert the information of the **city** into **area** since cities are a generalization of areas. Then we grab the next possible value in the subarea's *id* using a sequence and we add an instance of a **subarea** which will be used for the **Gym**. Again, **gym** represent a generalization of **subarea** so we need to add it there first. Finally, we repeat this process but for the **trainer** and **gym_leader** since **gym_leader** is a generalization from **trainer**.

8.7.3 Query validation

The trigger should have an effect everytime we look to insert a new **city**. Since we need to make sure that there is an area associated to the **city**, the trigger will take place before the insert. We can validate the trigger by doing the following insert:

```
INSERT INTO city (id, population) VALUES (400, 100000);
```

If we check all the information from **city** where the *id* is 400. We get the following information.

	id	population
	1	400

If we check the information of the **area** where the *id* is 400.

	id	name	name_region
	1	400 City	561 Lasalia

If we check the information about **subarea** where the *id_area* is 400.

	id	name	id_area
	1	651 566 Gym	400

If we check the information of the **gym_leader** whose gym is 651.

	id	gym
	1	537

Finally, if we see the information of the **trainer** whose id is 537.

	id	class	name	experience	gold
1	537	Gym Leader	Gym Leader 387	942	9461

8.8 Trigger 2

8.8.1 Solution

```

CREATE TABLE IF NOT EXISTS Warnings (
    id SERIAL PRIMARY KEY,
    trainer_id INT,
    species_name VARCHAR(255),
    route_name VARCHAR(255),
    message VARCHAR(255),
    date_obtained DATE
);

CREATE OR REPLACE FUNCTION check_new_pokemon()
RETURNS TRIGGER AS $$
DECLARE
    pavement_type VARCHAR(255);
    pokemon_min_level INT;
    pokemon_max_level INT;
    capture_chance INT;
    trainer_name VARCHAR(255);
    species_name VARCHAR(255);
    route_name VARCHAR(255);
BEGIN
    SELECT r.pavement INTO pavement_type
    FROM route r JOIN subarea s ON s.id_area = r.id
    WHERE s.id = NEW.subarea;

    SELECT p.name INTO species_name
    FROM species p WHERE p.id = NEW.species;

    SELECT t.name INTO trainer_name
    FROM trainer t WHERE t.id = NEW.trainer;

    SELECT f.min_level, f.max_level, f.probability
    INTO pokemon_min_level, pokemon_max_level, capture_chance
    FROM find f
    JOIN encounter e ON f.id_encounter = e.id
    WHERE f.id_species = NEW.species AND e.subarea = NEW.subarea
    LIMIT 1;

    SELECT a.name INTO route_name FROM area a
    JOIN subarea s ON a.id = s.id_area WHERE s.id = NEW.subarea;

    IF (pavement_type = 'water' AND NEW.method <> 'surf') THEN
        INSERT INTO Warnings (trainer_id, species_name, route_name, message,
        date_obtained)
        VALUES (NEW.trainer, species_name, route_name,
        CONCAT('Trainer ', trainer_name, ' captured ', species_name,
        ' in ', route_name,
        ' with an incorrect method, level, or capture chance.')),
        NEW.date_time;
    ELSIF (pavement_type <> 'water' AND (NEW.method <> 'walk' AND NEW.method <>
    'headbutt')) THEN
        INSERT INTO Warnings (trainer_id, species_name, route_name, message,
        date_obtained)
        VALUES (NEW.trainer, species_name, route_name,
        'The method chosen is invalid for this route type.');
    END IF;
END;
$$

```

```

        CONCAT('Trainer ', trainer_name, ' captured ', species_name,
        ' in ', route_name,
        ' with an incorrect method, level, or capture chance.'),
NEW.date_time);
ELSIF (NEW.level < pokemon_min_level OR NEW.level > pokemon_max_level OR
capture_chance <= 0) THEN
    INSERT INTO Warnings (trainer_id, species_name, route_name, message,
date_obtained)
VALUES (NEW.trainer, species_name, route_name,
CONCAT('Trainer ', trainer_name, ' captured ', species_name,
' in ', route_name,
' with an incorrect method, level, or capture chance.'),
NEW.date_time);
END IF;

RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER check_new_pokemon_trigger
AFTER INSERT ON pokemon_creature
FOR EACH ROW
EXECUTE FUNCTION check_new_pokemon();

```

8.8.2 Explanation

Firstly, we need to create the table **Warnings** which will contain the corresponding information of the wrong insert into the **pokemon_creature** table. We then do a series of queries to obtain the information regarding the pavement of the **route** where the pokemon was obtained, the name of its **species**, the name of the **trainer**, and the min level, *max level* and chance of it being obtained in that **subarea**. All of these queries are performed with the corresponding information of the newly obtained pokemon. After this, we do the corresponding checks and if at any moment the information inserted does not correspond with the conditions stated then we add a new warning. The conditions include, that if it was captured in water it has to have happened while surfing, if it isn't water it should have been with the headbutt or walking methods and the level should be between the min and max levels and with a probability higher than 0%.

8.8.3 Query validation

If we perform the following insertion:

```

INSERT INTO pokemon_creature (id, species, nickname, trainer, level, experience,
gender, nature, date_time, subarea, method, pokeball, is_baby, position, hp,
status_condition, team, attack, defense, special_attack, special_defense, speed,
evasion)
VALUES (13713, 16, 'bird', 24, 23, 5423, 'male', 23, '2021-08-29
04:24:54.000000', 22, 'surf', 3, false, null, 70, null, null,
100,100,100,100,100);

```

We get the following result in the **Warnings** table:

	id	trainer_id	species_name	route_name	message	date_obtained
1	1	24	pidgey	route 15	Trainer Wally captured pidgey in route 15 with an inco...	2021-08-29

In here we added a pokemon caught in a route which pavement is dirt but we use the surf method so that its why it has been added to the warnings table.

	id	name		id_area
1	22	route 15 area		285

	id	pavement	
1	285	Dirt	

If we do another test:

```
INSERT INTO pokemon_creature (id, species, nickname, trainer, level, experience, gender, nature, date_time, subarea, method, pokeball, is_baby, position, hp, status_condition, team, attack, defense, special_attack, special_defense, speed, evasion)
VALUES (13715, 98, 'crab', 24, 23, 5423, 'male', 23, '2021-08-29 04:24:54.000000', 193, 'walk', 3, false, null, 70, null, null, 100, 100, 100, 100, 100);
```

We get the following result in the **Warnings** table:

	r_id	species_name	route_name	message
1	24	pidgey	route 15	Trainer Wally captured pidgey in route 15 with an incorrect method, level, or capture chance.
2	24	krabby	route 40	Trainer Wally captured krabby in route 40 with an incorrect method, level, or capture chance.

In this case the pavement is water but we have used the method walk.

	id	name		id_area
1	193	route 40 area		22

	id	pavement	
1	22	Water	

And if we do another test:

```
INSERT INTO pokemon_creature (id, species, nickname, trainer, level, experience, gender, nature, date_time, subarea, method, pokeball, is_baby, position, hp, status_condition, team, attack, defense, special_attack, special_defense, speed, evasion)
VALUES (13716, 72, 'octopus', 24, 25, 5423, 'male', 23, '2021-08-29 04:24:54.000000', 193, 'surf', 3, false, null, 70, null, null, 100, 100, 100, 100, 100);
```

We get this result:

	id	trainer_id	species_name	route_name	message	date_obtained
1	1	24	pidgey	route 15	Trainer Wally captured pidgey in route 15 with an inco...	2021-08-29
2	2	24	krabby	route 40	Trainer Wally captured krabby in route 40 with an inco...	2021-08-29
3	3	24	tentacool	route 40	Trainer Wally captured tentacool in route 40 with an i...	2021-08-29

In this case it is the same area as before so it satisfies the condition of it being water and the method being surf but the level surpasses the max level.

	<code>id</code>	<code>id_encounter</code>	<code>id_species</code>	<code>condition_type</code>	<code>condition_value</code>	<code>probability</code>	<code>min_level</code>	<code>max_level</code>
1	3526	536	72			60	20	24
2	3527	536	72			30	15	19
3	3706	536	73			10	20	24

9 Cross queries.

9.1 Query 1

9.1.1 Solution

```
SELECT r.name
  FROM region r JOIN criminal_organization co ON r.name = co.region
  JOIN evil_criminal ec ON co.name = ec.criminal_organization
  JOIN team tc ON tc.trainer = ec.id JOIN battle b ON tc.id = b.winner
  JOIN team t ON t.id = b.loser
 GROUP BY r.name
 ORDER BY COUNT(t.id) DESC
 LIMIT 1;
```

	name
1	hoenn

9.1.2 Explanation

To find out the region that is most affected by a criminal region we start by joining the tables **Region** and **Criminal_Organization** (linking the *region* attribute of the criminal organization to the region's id). Then, to find all the battles that criminals have won we join the **Evil_Criminal's Team** with the *winner* attribute of the **Battle** entity (by means of two *JOINS*). Once we have these battles, we link them with the **Teams** who lost (*loser* attribute of the **Battle** table). This gives us the list of all the trainers defeated by criminals. To count how many happened in each area, we use a *GROUP BY* clause and merge all the rows of trainers defeated in the same region. Then, by means of the *COUNT()* operator, we are able to check how many defeated trainers each region had. Lastly, to show only the one that had the most criminal activity, we use the *ORDER BY* command to sort them by the number of defeated trainers and use the *LIMIT 1* operation to get only one row.

9.1.3 Query validation

Fristly, if we remove the *LIMIT* on the previous query and also select the *COUNT(t.id)* we use to sort the result, we can check the number of defeated trainers in each region:

	name	count
1	hoenn	53
2	johto	33
3	kanto	21
4	sinnoh	12

As you can see, we only get four regions in which criminal members defeated other trainers. To check if this makes sense, we can look at the **Criminal_Organization** table and count how many different regions they can be found in. As you can see below, there are only four regions in which the organizations are located. Therefore, the result of our query is consistent.

	name	building	base_location	region
1	Plasma	N's Castle		3 kanto
2	Rocket	Silph Co.		89 kanto
3	Galactic	Galactic Veilstone Building		142 sinnoh
4	Magma	Magma Hideout		224 hoenn
5	Aqua	Aqua Hideout		224 hoenn
6	Cipher	Souvenir Shop		351 johto

We will now add a new battle in which a criminal from an organization located at ‘johto’ region defeats another trainer and see how this affects the output of our query. As you can see in the table above, the ‘Cipher’ criminal organization is the one located in the ‘johto’ region. Below you can see that one of the evil criminals that belongs to this organization is the one with $id = 434$.

	id	money_steal	sig...	is_leader	sidekick	criminal_organization
	434	139	I mean, sa...	false		427 Cipher

We will execute the following query to insert a new battle in which the criminal with $id = 434$ defeats the trainer with $id = 0$ by first checking their respective teams in the **Team** table:

	id	trainer		id	trainer
	196	0		527	434

```
INSERT INTO battle(id, winner, loser) VALUES (700, 527, 196);
```

If we execute the query again, we get the following output:

	name	count
	hoenn	53
	johto	34
	kanto	21
	sinnoh	12

As you can see, the ‘johto’ region now has increased its defeated trainers by one unit. Meaning that the way we count the defeated trainers of each region is valid.

9.2 Query 2

9.2.1 Solution

```
SELECT pc.id, s.name AS species_name, a.name AS area_name
FROM pokemon_creature pc
JOIN species s ON pc.species = s.id
JOIN subarea sa ON pc.subarea = sa.id
JOIN area a ON sa.id_area = a.id
WHERE sa.id NOT IN (SELECT sa2.id FROM subarea sa2
    JOIN encounter e ON pc.subarea = e.subarea
    JOIN find f ON e.id = f.id_encounter
    JOIN species s2 ON f.id_species = s2.id
    WHERE s2.id = s.id);
```

9.2.2 Explanation

The essential logic behind the query is that we want to find those pokemon that have been captured in an **area** where it would have been impossible to find them there. Basically, in the subquery we find all the possible **subareas** where each pokemon could be found. Once we have that, we try and find all the pokemon caught in a **subarea** that did not exist in the possibilities of the **encounters**. We basically compare if the **subarea** where the pokemon was caught with the **subareas** that exist in the **encounters**. We took this approach since we noticed that there are no null values in the **probability** column.

9.2.3 Query validation

Initially, the query returned no results. However once we perform a insert we can see that the query works.

```
INSERT INTO pokemon_creature (id, species, nickname, trainer, level, experience, gender, nature, date_time, subarea, method, pokeball, is_baby, position, hp, status_condition, team, attack, defense, special_attack, special_defense, speed, evasion)
VALUES (13760, 460, 'Pedro', 1, 40, 150664, 'male', 18, '2022-02-01 07:58:51.000000', 41, 'walk', 450, false, null, 506, null, null, 50, 50, 50, 50, 50);
```

In this case, we are inserting the pokemon abomasnow that has been caught in subarea 41 when its only possible subareas where it can be caught are 385 and 386.

	id	species_name	area_name
1	13760	abomasnow	route 5

If we were to change the area where it was caught from subarea 41 to 385 then it will not be included in the results.

9.3 Query 3

9.3.1 Solution

9.3.2 Explanation

9.3.3 Query validation

We faced some problems when doing this trigger and didn't have enough time to fix it.

9.4 Query 4

9.4.1 Solution

```
SELECT tr.name AS gym_leader, gl.gym AS gym, t.name AS gym_type,
       p.id AS pokemon_id, s.name AS species_name, ht.type_name AS pokemon_type,
       tm.id AS team_id, m.name AS move_name, m.type AS move_type
FROM gym_Leader gl
JOIN gym g ON gl.gym = g.id
JOIN trainer tr ON gl.id = tr.id
JOIN team tm ON tr.id = tm.trainer
JOIN pokemon_Creature p ON tm.id = p.team
JOIN species s ON p.species = s.id
```

```

JOIN has_type ht ON p.species = ht.id_species
JOIN type t ON g.id_type = t.id
JOIN cause_or_receive cr ON p.id = cr.creature
JOIN move m ON cr.move = m.id
WHERE NOT EXISTS (
    SELECT 1
    FROM move mv
    WHERE mv.type = t.name
    AND mv.id = m.id
)
AND NOT EXISTS (
    SELECT 1
    FROM has_type ht2
    WHERE ht2.id_species = p.species
    AND ht2.type_name = t.name
)
ORDER BY gl.id, p.id, m.id;

```

The screenshot shows a table with 20 rows of data. The columns are labeled: gym_leader, gym, gym_type, pokemon_id, species_name, pokemon_type, team_id, move_name, and move_type. The data represents moves used by Brock's team members (species: spearow, hariyama, tangrowth, bronzer) against various opponents (team_id: 305). The moves listed are take down, peck, smokescreen, false swipe, karate chop, rolling kick, detect, rock smash, mega drain, solar beam, petal dance, spore, confusion, light screen, and light screen again.

	gym_leader	gym	gym_type	pokemon_id	species_name	pokemon_type	team_id	move_name	move_type
1	Brock	543	rock	10331	spearow	flying	305	take down	normal
2	Brock	543	rock	10331	spearow	normal	305	take down	normal
3	Brock	543	rock	10331	spearow	flying	305	peck	flying
4	Brock	543	rock	10331	spearow	normal	305	peck	flying
5	Brock	543	rock	10331	spearow	flying	305	smokescreen	normal
6	Brock	543	rock	10331	spearow	normal	305	smokescreen	normal
7	Brock	543	rock	10331	spearow	flying	305	false swipe	normal
8	Brock	543	rock	10331	spearow	normal	305	false swipe	normal
9	Brock	543	rock	10332	hariyama	fighting	305	karate chop	fighting
10	Brock	543	rock	10332	hariyama	fighting	305	rolling kick	fighting
11	Brock	543	rock	10332	hariyama	fighting	305	detect	fighting
12	Brock	543	rock	10332	hariyama	fighting	305	rock smash	fighting
13	Brock	543	rock	10333	tangrowth	grass	305	mega drain	grass
14	Brock	543	rock	10333	tangrowth	grass	305	solar beam	grass
15	Brock	543	rock	10333	tangrowth	grass	305	petal dance	grass
16	Brock	543	rock	10333	tangrowth	grass	305	spore	grass
17	Brock	543	rock	10334	bronzer	steel	305	confusion	psychic
18	Brock	543	rock	10334	bronzer	psychic	305	confusion	psychic
19	Brock	543	rock	10334	bronzer	steel	305	light screen	psychic
20	Brock	543	rock	10334	bronzer	psychic	305	light screen	psychic

9.4.2 Explanation

In order to perform this query we need to relate a lot of different tables. We start at **gym_leader** and we go on to find their pokémon that belong to a **team** and which **moves** do they have. The central part of the query are the subqueries used in the WHERE clause. Both conditions in the first subquery aim to relate first the **species** with their **type** and then it makes sure that the **gym type** and the **pokémon type** match. With this we filter out rows where a match exists in the **has_type** table for the current pokémon **species** and the gym type and we end up displaying only those pokémon that do not share the **gym's type**. We do something similar in the second subquery where we aim to exclude every move that has the same **type** as the gym.

9.4.3 Query validation

For example, we can see that Blaine has a luvdisc in his team which is water type. But if we changed that luvdisc's type to fire then it will not appear:

	gym_leader	gym	gym_type	pokemon_id	species_name	pokemon_type	team_id	move_name	move_type
1	Blaine	557	fire	10426	luvdisc	water	46	hydro pump	water
2	Blaine	557	fire	10426	luvdisc	water	46	waterfall	water
3	Blaine	557	fire	10426	luvdisc	water	46	octazooka	water
4	Blaine	557	fire	10426	luvdisc	water	46	whirlpool	water
5	Blaine	557	fire	10427	slowking	psychic	46	hypnosis	psychic
6	Blaine	557	fire	10427	slowking	water	46	hypnosis	psychic
7	Blaine	557	fire	10427	slowking	psychic	46	agility	psychic
8	Blaine	557	fire	10427	slowking	water	46	agility	psychic
9	Blaine	557	fire	10427	slowking	psychic	46	rest	psychic
10	Blaine	557	fire	10427	slowking	water	46	rest	psychic
11	Blaine	557	fire	10427	slowking	water	46	mirror coat	psychic
12	Blaine	557	fire	10427	slowking	psychic	46	mirror coat	psychic
13	Blaine	557	fire	10428	dragonite	dragon	46	gust	flying
14	Blaine	557	fire	10428	dragonite	flying	46	gust	flying
15	Blaine	557	fire	10428	dragonite	dragon	46	peck	flying
16	Blaine	557	fire	10428	dragonite	flying	46	peck	flying
17	Blaine	557	fire	10428	dragonite	flying	46	aeroblast	flying
18	Blaine	557	fire	10428	dragonite	dragon	46	aeroblast	flying
19	Blaine	557	fire	10428	dragonite	flying	46	twister	dragon
20	Blaine	557	fire	10428	dragonite	dragon	46	twister	dragon

```
UPDATE has_type
SET type_name = 'fire'
WHERE id_species = 370;
```

	gym_leader	gym	gym_type	pokemon_id	species_name	pokemon_type	team_id	move_name	move_type
1	Blaine	557	fire	10427	slowking	water	46	hypnosis	psychic
2	Blaine	557	fire	10427	slowking	psychic	46	hypnosis	psychic
3	Blaine	557	fire	10427	slowking	psychic	46	agility	psychic
4	Blaine	557	fire	10427	slowking	water	46	agility	psychic
5	Blaine	557	fire	10427	slowking	water	46	rest	psychic
6	Blaine	557	fire	10427	slowking	psychic	46	rest	psychic
7	Blaine	557	fire	10427	slowking	psychic	46	mirror coat	psychic
8	Blaine	557	fire	10427	slowking	water	46	mirror coat	psychic
9	Blaine	557	fire	10428	dragonite	dragon	46	gust	flying
10	Blaine	557	fire	10428	dragonite	flying	46	gust	flying
11	Blaine	557	fire	10428	dragonite	flying	46	peck	flying
12	Blaine	557	fire	10428	dragonite	dragon	46	peck	flying
13	Blaine	557	fire	10428	dragonite	flying	46	aeroblast	flying
14	Blaine	557	fire	10428	dragonite	dragon	46	aeroblast	flying
15	Blaine	557	fire	10428	dragonite	flying	46	twister	dragon
16	Blaine	557	fire	10428	dragonite	dragon	46	twister	dragon
17	Blaine	557	fire	10429	eevee	normal	46	roar	normal
18	Blaine	557	fire	10429	eevee	normal	46	sonic boom	normal
19	Blaine	557	fire	10429	eevee	normal	46	screech	normal
20	Blaine	557	fire	10429	eevee	normal	46	enrage	normal

And if we change the type of a move:

```
UPDATE move
SET type = 'rock'
WHERE name = 'take down';
```

We can see that Brock's Spearow still appears but only three of its moves do:

	gym_leader	gym	gym_type	pokemon_id	species_name	pokémon_type	team_id	move_name	move_type
1	Brock	543	rock	10331	spearow	flying	305	peck	flying
2	Brock	543	rock	10331	spearow	normal	305	peck	flying
3	Brock	543	rock	10331	spearow	flying	305	smokescreen	normal
4	Brock	543	rock	10331	spearow	normal	305	smokescreen	normal
5	Brock	543	rock	10331	spearow	normal	305	false swipe	normal
6	Brock	543	rock	10331	spearow	flying	305	false swipe	normal
7	Brock	543	rock	10332	hariyama	fighting	305	karate chop	fighting
8	Brock	543	rock	10332	hariyama	fighting	305	rolling kick	fighting
9	Brock	543	rock	10332	hariyama	fighting	305	detect	fighting
10	Brock	543	rock	10332	hariyama	fighting	305	rock smash	fighting
11	Brock	543	rock	10333	tangrowth	grass	305	mega drain	grass
12	Brock	543	rock	10333	tangrowth	grass	305	solar beam	grass
13	Brock	543	rock	10333	tangrowth	grass	305	petal dance	grass
14	Brock	543	rock	10333	tangrowth	grass	305	spore	grass
15	Brock	543	rock	10334	bronzor	steel	305	confusion	psychic
16	Brock	543	rock	10334	bronzor	psychic	305	confusion	psychic
17	Brock	543	rock	10334	bronzor	steel	305	light screen	psychic
18	Brock	543	rock	10334	bronzor	psychic	305	light screen	psychic
19	Brock	543	rock	10334	bronzor	steel	305	amnesia	psychic
20	Brock	543	rock	10334	bronzor	psychic	305	amnesia	psychic

9.5 Query 5

9.5.1 Solution

```
SELECT ec.id AS criminal, pc.id AS pokémon
  FROM evil_criminal ec JOIN pokemonCreature pc ON pc.trainer = ec.id
    JOIN causeOrReceive cor ON pc.id = cor.creature
    JOIN move m ON cor.move = m.id JOIN healing h ON m.id = h.id
   WHERE pc.method = 'gift' AND ec.isLeader = TRUE
     AND m.accuracyPercentage != 0 AND pc.subarea IS NOT NULL;
```

criminal	pokémon

As you can see, the query does not return any rows. We will explain the reason why below.

9.5.2 Explanation

To get the Pokémons that have been obtained through a ‘gift’ by a criminal we first join these two tables (**Evil_Criminal** and **Pokemon_Creature**) and specify in the *WHERE* clause that the Pokémons’ *method* (representing how they were obtained) equals ‘gift’. To ensure now these were obtained in a subarea, we add a second condition to the *WHERE* clause specifying that the *subarea* attribute of the **Pokemon_Creature** table cannot be *NULL*. Also, to show only the criminal leaders we add another condition to the *WHERE* clause stating that their *is_leader* attribute has to be *TRUE*.

Lastly, to check if the Pokémons have a move that can drain other Pokémons HP we join the **Pokemon_Creature** and **Move** entities (through the **Cause_Or_Receive** table). Now, as we classified the drain moves as **Healing** moves on the previous phase (as the Pokémons use them to heal themselves) we join this table with the **Move**’s *id*. This will give us all the moves that are of type healing. However, we realized that the drain moves are those that have a none zero in their *accuracy_percentage* attribute (as opposed to all other healing moves that don’t drain). Hence, we added this condition to the *WHERE* clause.

9.5.3 Query validation

This query did not return any rows because there aren't any Pokémons in the given datasets that satisfy all these conditions. Therefore, to prove that the query works we will insert a new Pokémon that does by executing the query below. We will first add the Pokémon to the **Pokemon_Creature** table by adding a big *id* (not already added), giving it a *nickname*, allocate it a valid *subarea*, stating its obtention method as 'gift', and assign it to an evil criminal that is a leader (*id* = 507 as you can see on the screenshot below).

```
INSERT INTO pokemon_creature(id, species, nickname, subarea, method, trainer)
VALUES(14000, 1, 'ultimatePokemon', 127, 'gift', 507);
```

id	money_stole	sig...	is_leader
507	696970000000000	Hit the gy...	true

Also, we must ensure this Pokémon has a move that drains other Pokémons HP, so we create a new entry at the **Cause_Or_Receive** table linking this Pokémon to the move with *id* = 71, which as you can double check in the screenshot below is a draining move.

```
INSERT INTO cause_or_receive(creature, move, id) VALUES (14000, 71, 14000);
```

id	name	...	effect
71	absorb	100	Drains half the damage inflicted to heal the user.

Therefore, after inserting these two rows and executing the previous query again, we get the following:

criminal	pokémon
507	14000

As you can see, the query now does return the newly inserted Pokémon, meaning the conditions of the query are checked correctly, but before there weren't any Pokémon satisfying them.

Also, to prove that the drain moves are those we stored as *Healing* and have an accuracy percentage different than zero we will execute the following query:

```
SELECT m.name, m.effect
    FROM move m
    JOIN healing h on m.id = h.id
    WHERE accuracy_percentage != 0;
```

This is the output we get. Notice how all these are draining other Pokémons HP.

name	effect
absorb	Drains half the damage inflicted to heal the user.
mega drain	Drains half the damage inflicted to heal the user.
dream eater	Only works on sleeping Pokémon. Drains half the damage inflicted to heal the user.
leech life	Drains half the damage inflicted to heal the user.
giga drain	Drains half the damage inflicted to heal the user.
drain punch	Drains half the damage inflicted to heal the user.
horn leech	Drains half the damage inflicted to heal the user.
parabolic charge	Heals the user for half the total damage dealt to all targets.
draining kiss	Drains 75% of the damage inflicted to heal the user.
oblivion wing	Drains 75% of the damage inflicted to heal the user.

However, if we invert the *WHERE* condition of the previous query and make it equal to zero we obtain these:

name	effect
recover	Heals the user by half its max HP.
soft boiled	Heals the user by half its max HP.
milk drink	Heals the user by half its max HP.
morning sun	Heals the user by half its max HP. Affected by weather.
synthesis	Heals the user by half its max HP. Affected by weather.
moonlight	Heals the user by half its max HP. Affected by weather.
swallow	Recovers 1/4 HP after one Stockpile, 1/2 HP after two Stockpiles, or full HP after three Stockpiles.
slack off	Heals the user by half its max HP.
roost	Heals the user by half its max HP.
heal order	Heals the user by half its max HP.
heal pulse	Heals the target for half its max HP.
shore up	Heals the user for $\frac{1}{2}$ its max HP, or $\frac{2}{3}$ during a sandstorm.
floral healing	Heals the target for $\frac{1}{2}$ its max HP, or $\frac{2}{3}$ on Grassy Terrain.
purify	Cures the target of a major status ailment and heals the user for 50% of its max HP.
jungle healing	Inflicts regular damage with no additional effect.

Notice how all these aren't drain moves.

9.6 Query 6

9.6.1 Solution

9.6.2 Explanation

9.6.3 Query validation

We faced some problems when doing this trigger and didn't have enough time to fix it.

9.7 Query 7

9.7.1 Solution

```
SELECT p.id, p.species, p.nickname, p.level, p.experience,
p.nature,
      p.hp, (((2 * hp + (level ^ 2) / 4) / 100 + 5) * nature) as
update_hp,
      p.attack, (((2 * p.attack + (level ^ 2) / 4) / 100 + 5) *
nature) as updated_attack,
      p.defense, (((2 * p.defense + (level ^ 2) / 4) / 100 + 5) *
nature) as updated_defense,
```

```

        p.special_attack, (((2 * p.special_attack + (level ^ 2) / 4)
/ 100 + 5) * nature) as updated_special_attack,
        p.special_defense, (((2 * p.special_defense + (level ^ 2) /
4) / 100 + 5) * nature) as updated_special_defense,
        p.speed, (((2 * p.speed + (level ^ 2) / 4) / 100 + 5) *
nature) as updated_speed
FROM pokemon_creature as p
WHERE EXTRACT(YEAR FROM p.date_time) >= 2022 and EXTRACT(YEAR FROM
p.date_time) <= 2023;

```

Output Result 217

	id	species	nickname	level	experience	nature	hp	up
1	24	336	Flakey	17	3892	8	73	
2	94	88	Samantha	16	5261	11	80	
3	96	339	Tango	12	2378	14	50	
4	261	380	Bongo	1	31	8	80	
5	405	184	Jess	15	3384	2	100	
6	569	4	Ming	15	2754	6	39	
7	810	1	Napoleon	30	24801	14	45	114.1000000000000
8	819	44	Alex	31	28033	5	60	43.01249000000000
9	820	127	Dragster	35	62322	10	65	
10	1042	201	Tobie	25	19210	25	48	
11	1130	34	Eifel	27	18780	13	81	109.7524900000000
12	1133	292	Tippy	28	35116	4	1	
13	1212	116	Isabella	32	38709	1	30	
14	1224	57	Autumn	55	182993	17	65	235.6625000000000

Output Result 217

	hp	update_hp	attack	updated_attack	defense	updated_defense
1	73	57.46	100	61.78	60	55.38
2	80	79.64	80	79.64	50	73.03999999999999
3	50	89.04	48	88.48	43	87.08
4	80	52.82	80	52.82	90	54.42
5	100	15.125	50	13.125	80	14.325
6	39	38.055	52	39.615	43	38.535000000000004
7	45	114.1000000000001	49	115.22	49	115.22
8	60	43.01249999999996	65	43.5125	70	44.0125
9	65	93.625	125	105.625	100	100.625
10	48	188.0625	72	200.0625	48	188.0625
11	81	109.7524999999998	102	115.2125	77	108.7125
12	1	27.92	90	35.04	45	31.439999999999998
13	30	8.16	40	8.36	70	8.96
14	65	235.6625000000002	105	249.2625	60	233.9624999999998

	special_attack	updated_special_attack	special_defense	updated_special_defense
1	100	61.78	60	59.51
2	40	70.83999999999999	50	73.03999999999999
3	46	87.92	41	86.46
4	110	57.62000000000005	130	61.13
5	60	13.525	80	14.0
6	60	40.575	50	39.5
7	65	119.7000000000002	65	119.7000000000001
8	85	45.51249999999996	75	44.95
9	55	91.625	70	94.0
10	72	200.0625	48	188.0
11	85	110.79250000000002	75	108.0
12	30	30.24000000000002	30	30.24000000000001
13	70	8.96	25	8.76
14	60	233.96249999999998	70	237.3625

	special_attack	special_defense	updated_special_defense	speed	updated_speed
1	61.78	60	55.38	65	56.18
2	39999999999999	50	73.03999999999999	25	67.53999999999999
3	87.92	41	86.52	60	91.84
4	10000000000005	130	60.82	110	57.62000000000005
5	13.525	80	14.325	50	13.125
6	40.575	50	39.375	65	41.175
7	10000000000002	65	119.7000000000002	45	114.1000000000001
8	2499999999996	75	44.5125	40	41.0125
9	91.625	70	94.625	85	97.625
10	200.0625	48	188.0625	48	188.0625
11	9250000000002	75	108.1925	85	110.7925000000002
12	10000000000002	30	30.24000000000002	40	31.04
13	8.96	25	8.06	60	8.76
14	1624999999998	70	237.3625	95	245.8625

9.7.2 Explanation

This query simply looks for the `pokemon_creature` captured between 2022 and 2023 with the `WHERE` clause and `extract` function from PostgreSQL to take the year from the `datetime` attribute. Then, we simply took the same formula and put the updated statistic.

9.7.3 Query validation

To validate it, we have simply printed the level, nature and the statistics. If we apply the formula for hp 73, level 17 and nature 8, we will get approximately 57.46. Since this query has the same logic applied to each statistic, we can know that the rest of the statistics is fine as well.

9.8 Query 8

9.8.1 Solution

9.8.2 Explanation

9.8.3 Query validation

We faced some problems when doing this trigger and didn't have enough time to fix it.

10 Conclusions

10.1 Use of resources

Stage	Valèria	Eugènia	Sebastián	Marc	Total
ER/RM model update	0h	0h	0h	¼ h	¼ h
Data type selection	0h	0h	0h	0h	0h
Physical model codification	0h	0h	0h	0h	0h
Database population	0h	0h	0h	0 h	0 h
Implementation and validation of the queries and triggers	10h	10h	10h	10h	40h
Reporting	10h	7h	7h	7h	28h
Total:	20 h	20 h	20 h	20.25 h	80.25h

In general, the parts that took us longer to complete were the implementation of the queries and triggers, as we got stuck with some of them. Also, reporting everything that we did and explaining the validation of our queries was a long process.

We believed we all worked approximately the same amount of ours as we distributed the work equally between all the team members and helped one another when any of us encountered an issue they couldn't solve by themselves.

10.2 IA use (if required, 1-2 pages)

We did not use any AI tools to develop this project.

10.3 Lessons learnt and conclusions (1 page)

This stage of the project has allowed us to put into practice everything we have learnt throughout the course. Firstly, we were able to code hard queries in SQL that required a large set of JOINS and/or subqueries. In fact, it has forced us to review concepts that we had forgotten a bit such as the different kinds of joins or subqueries.

Secondly, we have had to review triggers as it was a topic we did very briefly in the first semester and we had happened to forget it by now. However, reviewing the slides made us remember the previously learned concepts and we were able to successfully implement almost all the required triggers.

Lastly, we have learned the importance of organization in a team in order to successfully carry out a project of this size. It is true that we already had to do it in the previous phase, but this one required more coordination between all the members of the team because of the little time we had to implement this one, considering the enormous amount of work we all have to do at the end of the semester. If we had had more time, we believe we could have finished all the work, as we faced some problems with some queries and didn't have time to fix them.

To conclude, we can say that this project has helped us to reinforce our SQL skills and taught us how to operate a database from its creation to its population and usage. Moreover, it has allowed us to have a satisfactory first contact with database programming.